

# 3011979 Intro to Deep Learning for Medical Imaging

## L11: Neural network architectures (for imaging data)

Apr 16<sup>th</sup>, 2021

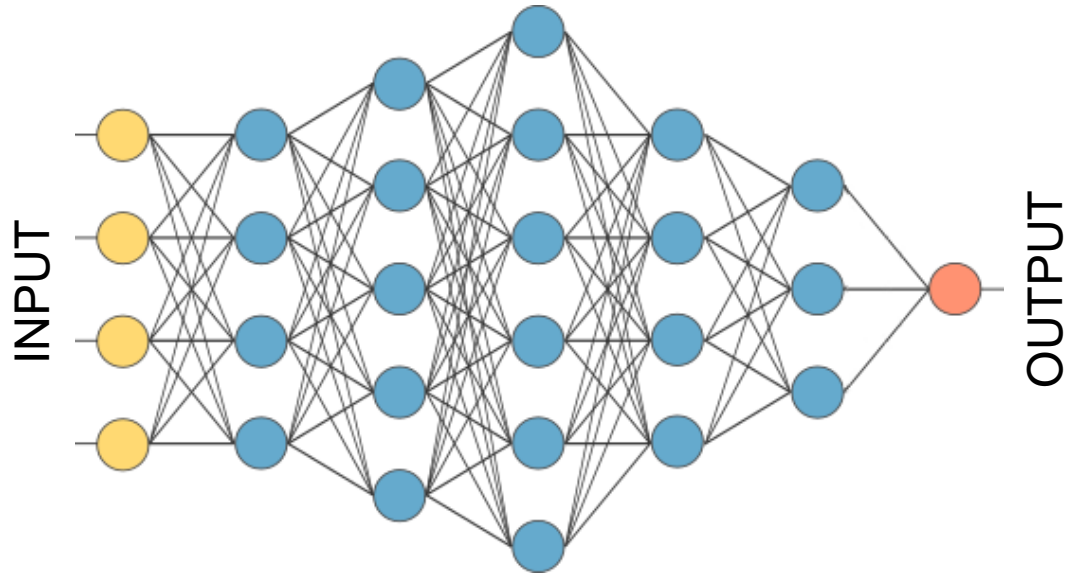


Sira Sriswasdi, Ph.D.

Research Affairs, Faculty of Medicine  
Chulalongkorn University

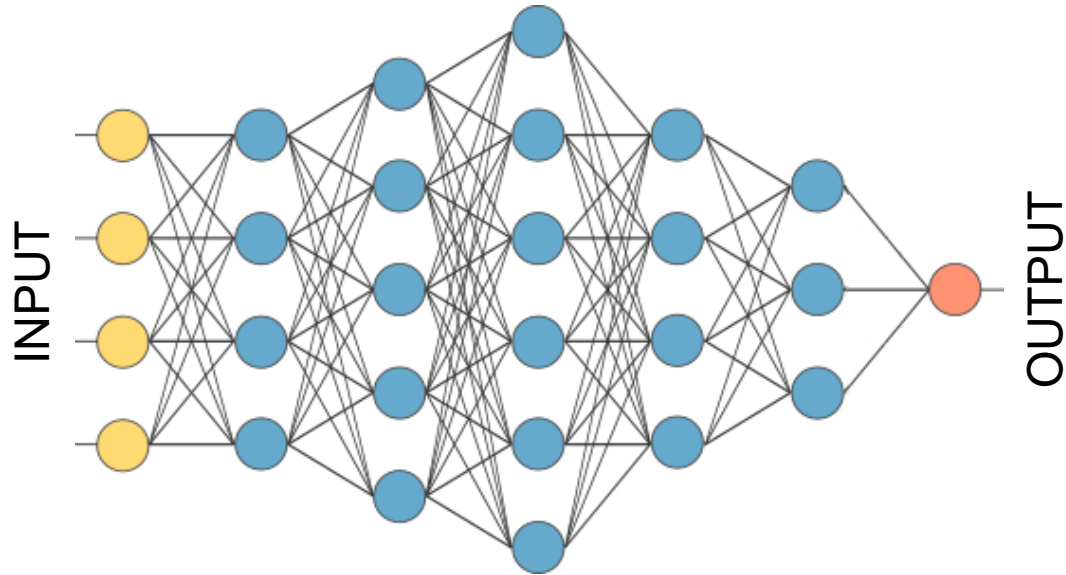
# Limitations of multilayer perceptron (fully connected / dense network)

# Fully connected network – powers



- Generate powerful predictor as a complex function of input features
- Can approximate any real-value function with simple, non-polynomial activation function

# Fully connected network – limitations



- What if features have contextual relationships?
  - Adjacent pixels in an image
  - Adjacent words in a sentence
- What if input size is variable?
  - Number of words in a medical report

# Fully connected network's image perception

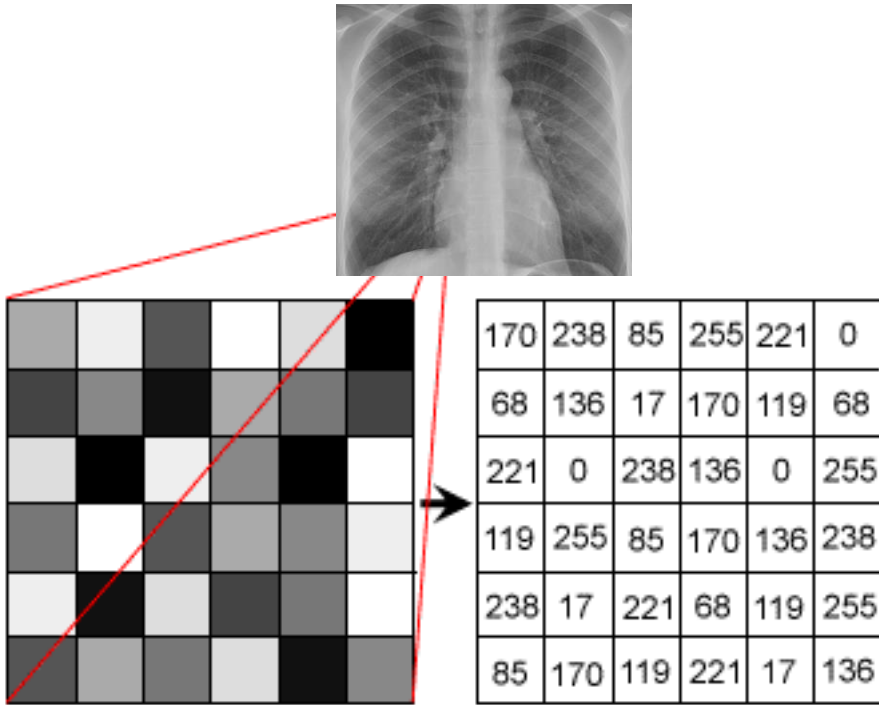
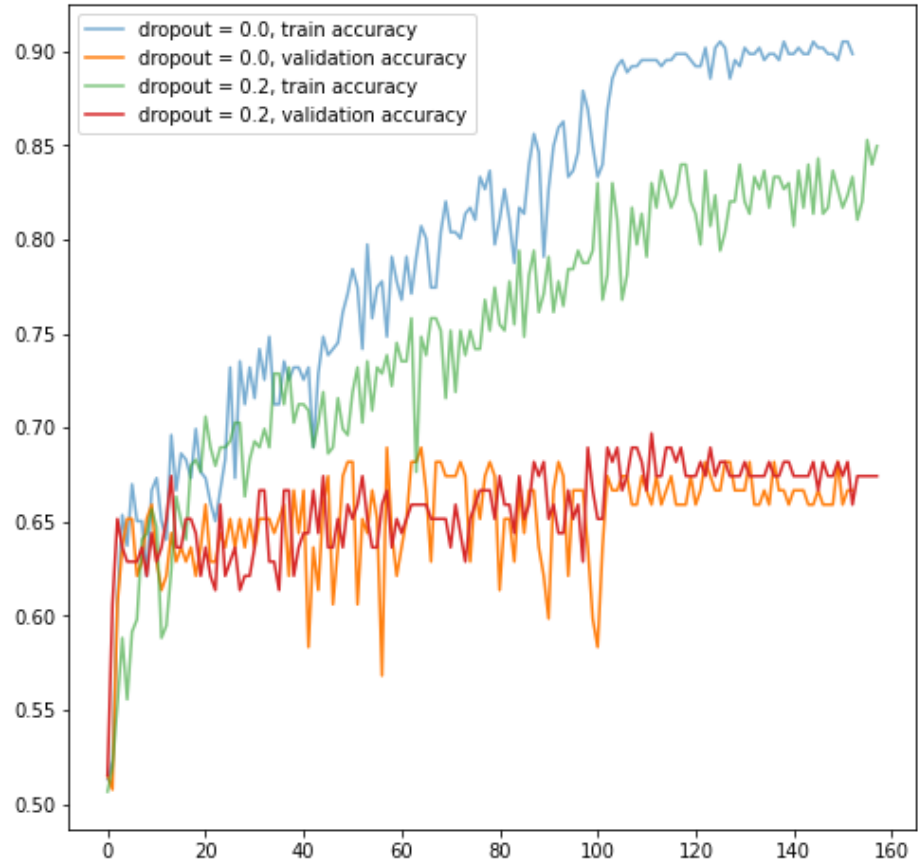


Image from [naushardsblog.wordpress.com](http://naushardsblog.wordpress.com)



- Fully connected network cannot learn anything from pixel values of chest x-ray images

# Image object

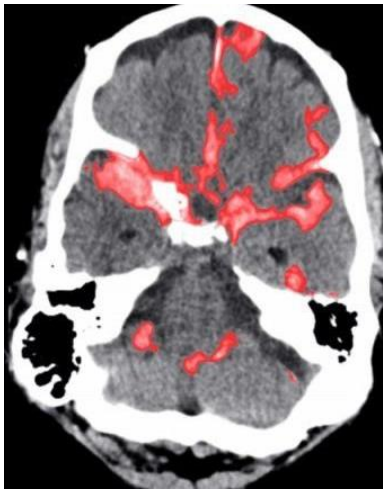
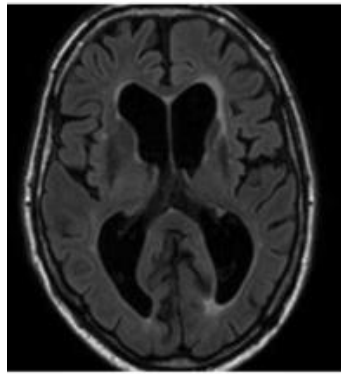
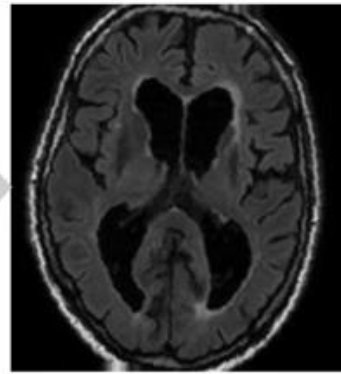


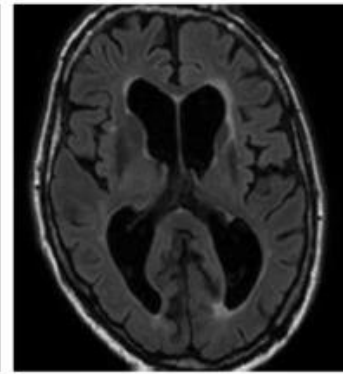
Image from tectales.com



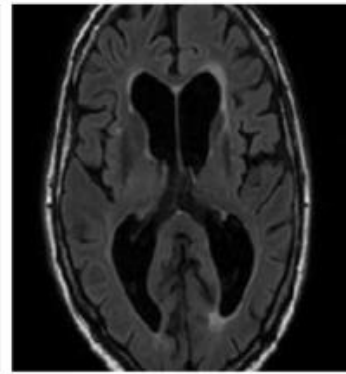
original slice



rotation



shear mapping



scaling

Li et al. Neuroimage 183 (2018)

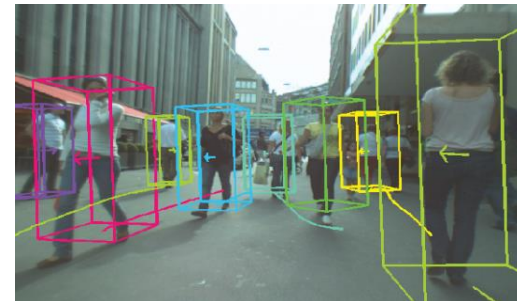
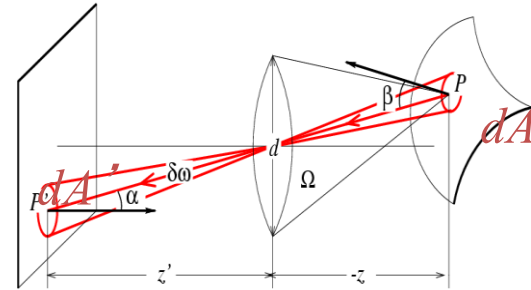
- An image object is a collection of related pixels
- Remain the same when rotated or translated
  - Fully connect network will perceive these variations as totally different inputs

# Computer vision

(Credit: Minh Hoai Nguyen, MLRS2019)

# Sciences of image

- Image processing
  - Image to image
- Imaging
  - Physics to image
- Graphics
  - Symbols to image
- Computer vision
  - Image to symbols



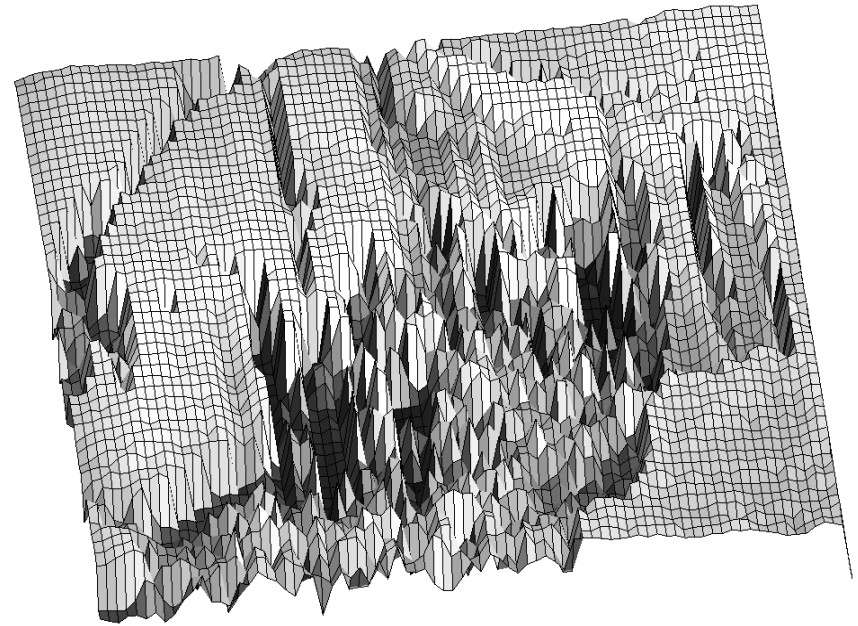


# Challenges of CV

Image through human eyes

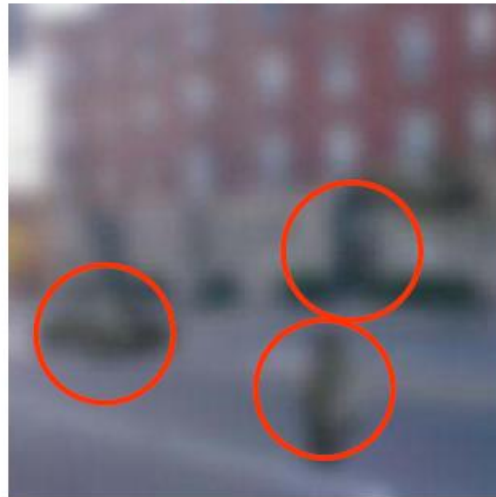


Image in computer

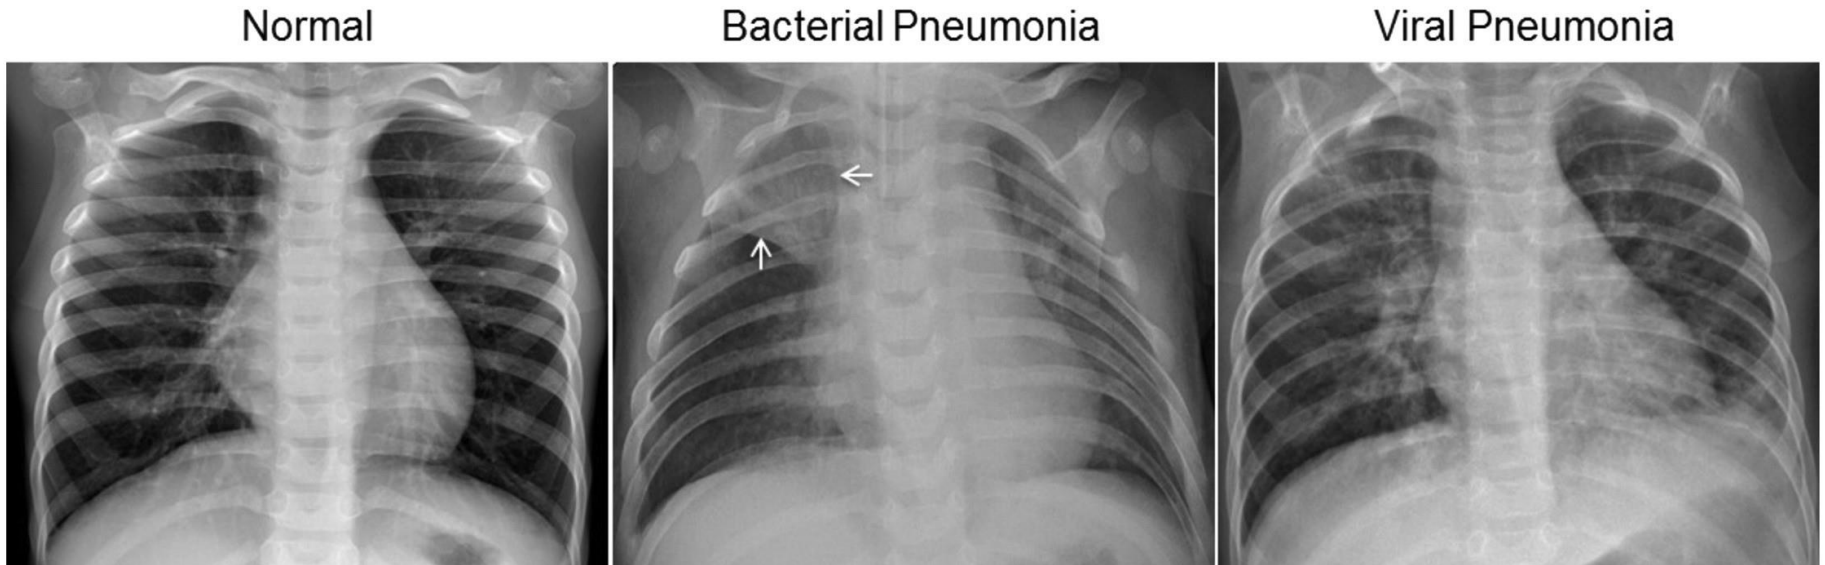


- We see more than just colors and numbers
- Try describing the concept “ugly” to a blind man
  - Or to a computer

# Context matters a lot



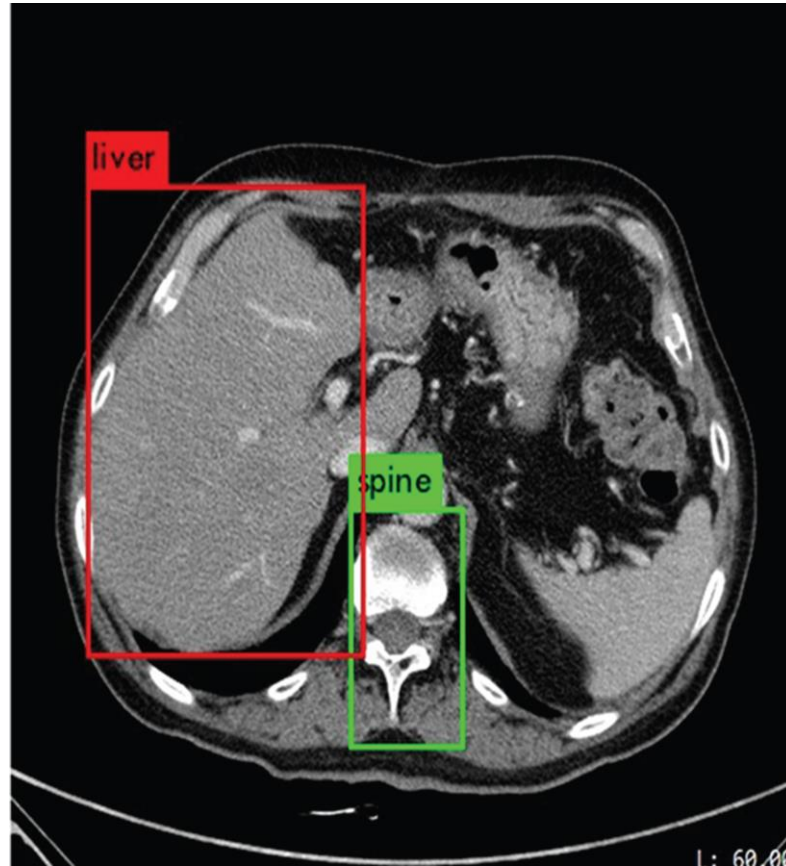
# Main tasks in CV – image classification



Keremany et al. Cell (2018)

- Input: An image
- Output: Single-label or multi-label classification

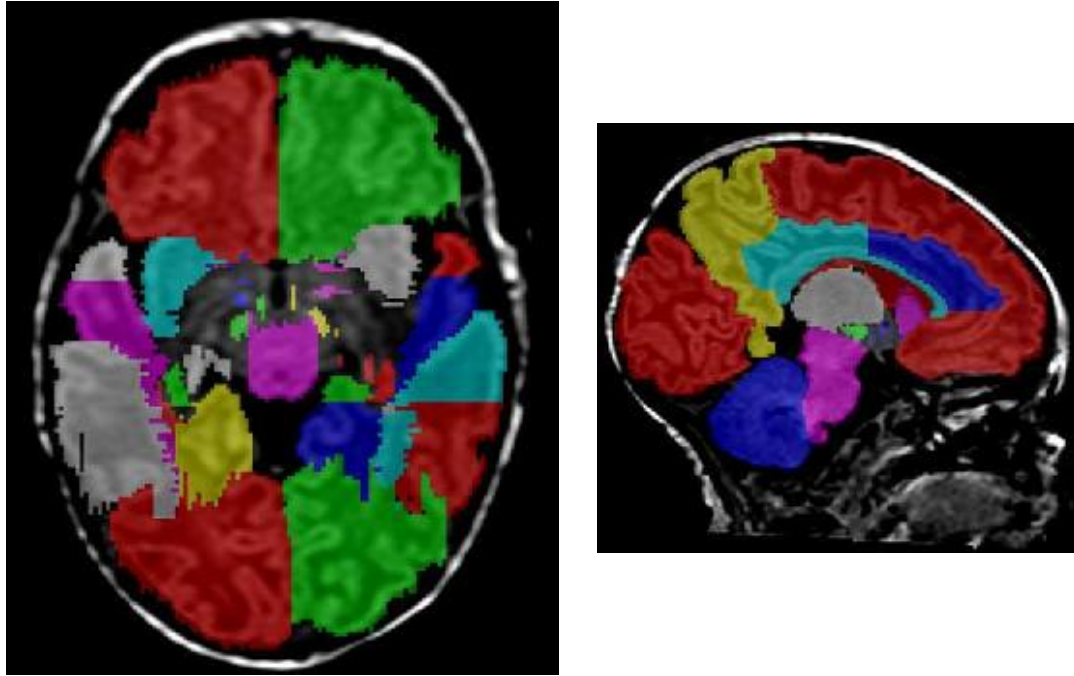
# Main tasks in CV – object detection



Pang et al. PLoS ONE (2019)

- Input: An image
- Output: Bounding box + Object label

# Main tasks in CV – semantic segmentation



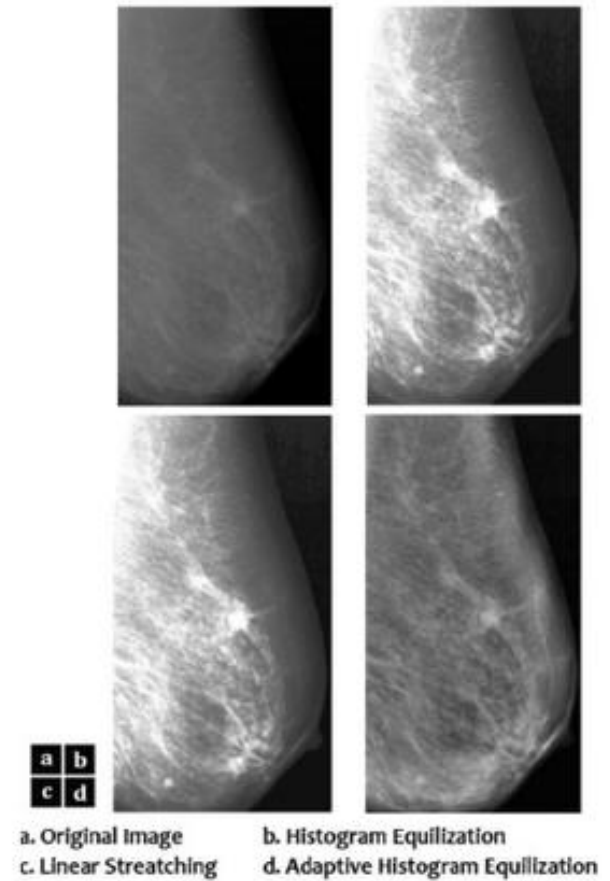
Makropoulos et al. IEEE Trans Med Imaging (2014)

- Input: An image
- Output: Pixel-level classification (and more)

# Image processing

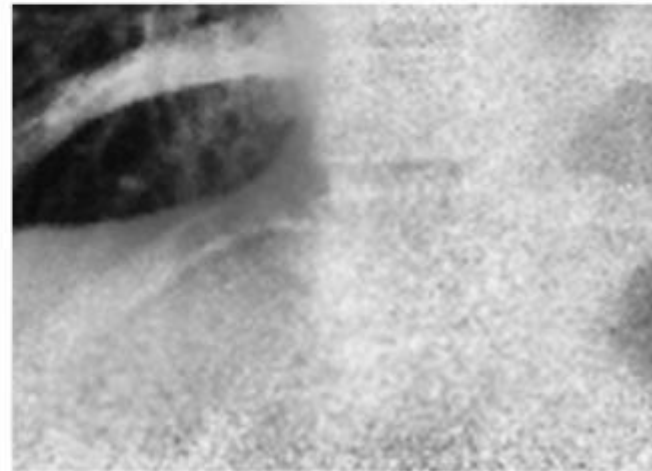
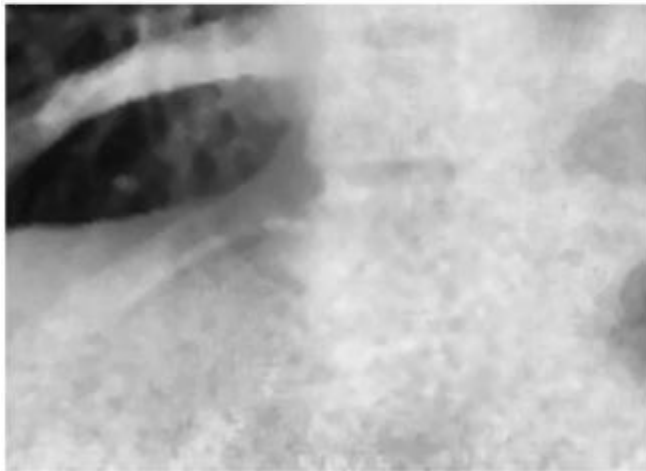
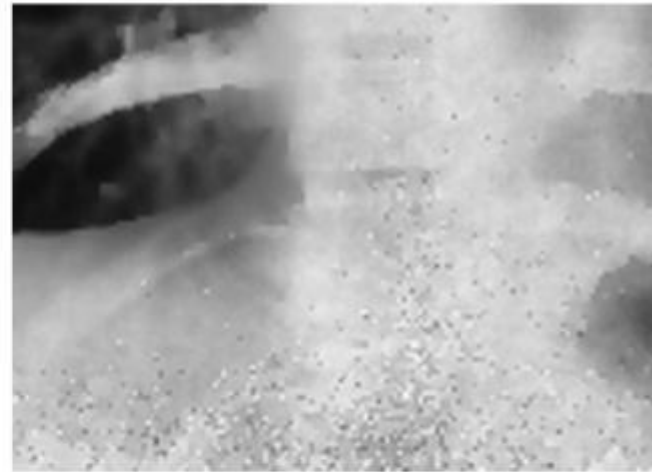
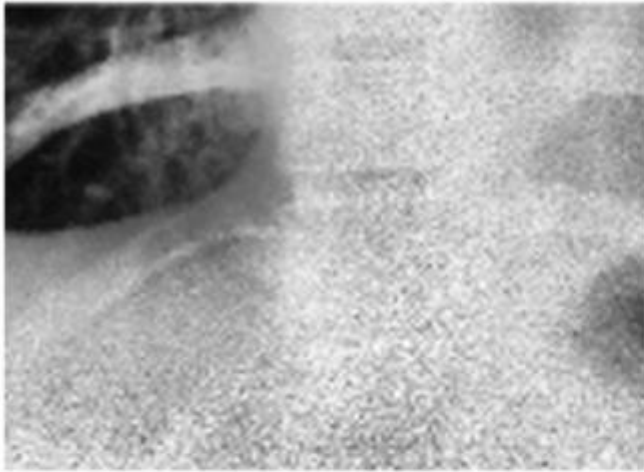


# Contrast adjustment



Ahmed et al. TechnoMoot conference (2011)

# Noise filtering



**a. Cropped CLAHE image**  
**c. Median Filter**

**b. Sigma Filter**  
**d. Wiener Filter**



# Mean filtering operation

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

	0								

- Setting center pixel = average of 3x3 area

# Mean filtering operation

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

	0	10							

- Setting center pixel = average of 3x3 area

# Mean filtering operation

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

	0	10	20						

- Setting center pixel = average of 3x3 area

# Mean filtering operation

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

- Output is a smoothed-out version of the input image

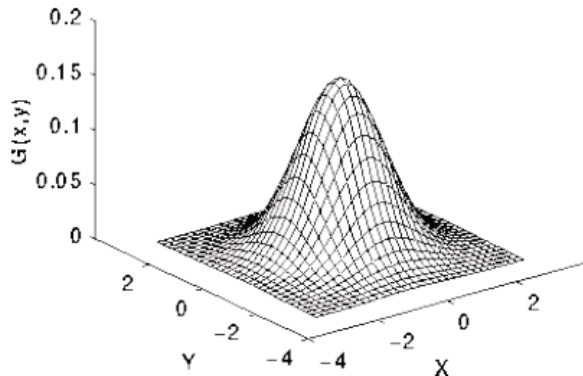
# Filtering as mathematical operation

- 3x3 mean filtering
- 5x5 Gaussian filtering
- Integrate contextual info!

$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$
$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$
$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$

$$G(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

$$\sigma = 1$$



$$\frac{1}{115}$$

2	4	5	4	2
4	9	12	9	4
5	12	15	12	5
4	9	12	9	4
2	4	5	4	2

# Filtering effects



input image

<https://setosa.io/ev/image-kernels/>

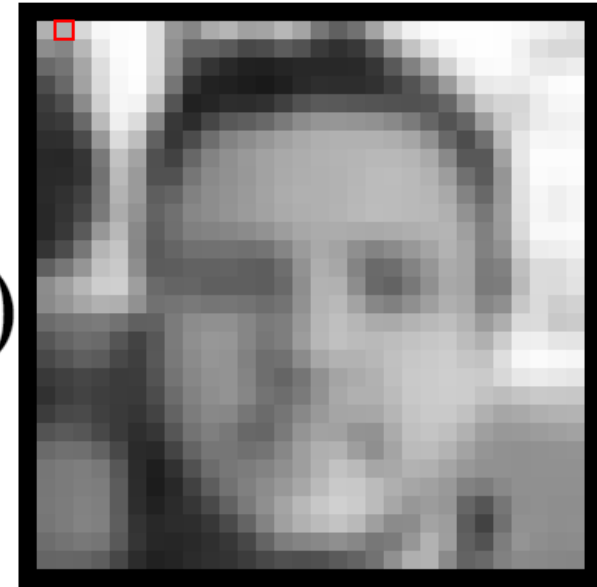


$$\begin{pmatrix} 205 + 247 + 245 \\ \times 0.0625 \times 0.125 \times 0.0625 \\ + 161 + 137 + 244 \\ \times 0.125 \times 0.25 \times 0.125 \\ + 154 + 75 + 200 \\ \times 0.0625 \times 0.125 \times 0.0625 \end{pmatrix}$$

$$= 175$$

kernel:

blur



output image

$$\begin{pmatrix} 205 + 247 + 245 \\ \times -1 \times -1 \times -1 \\ + 161 + 137 + 244 \\ \times -1 \times 8 \times -1 \\ + 154 + 75 + 200 \\ \times -1 \times -1 \times -1 \end{pmatrix}$$

$$= -435$$

kernel:

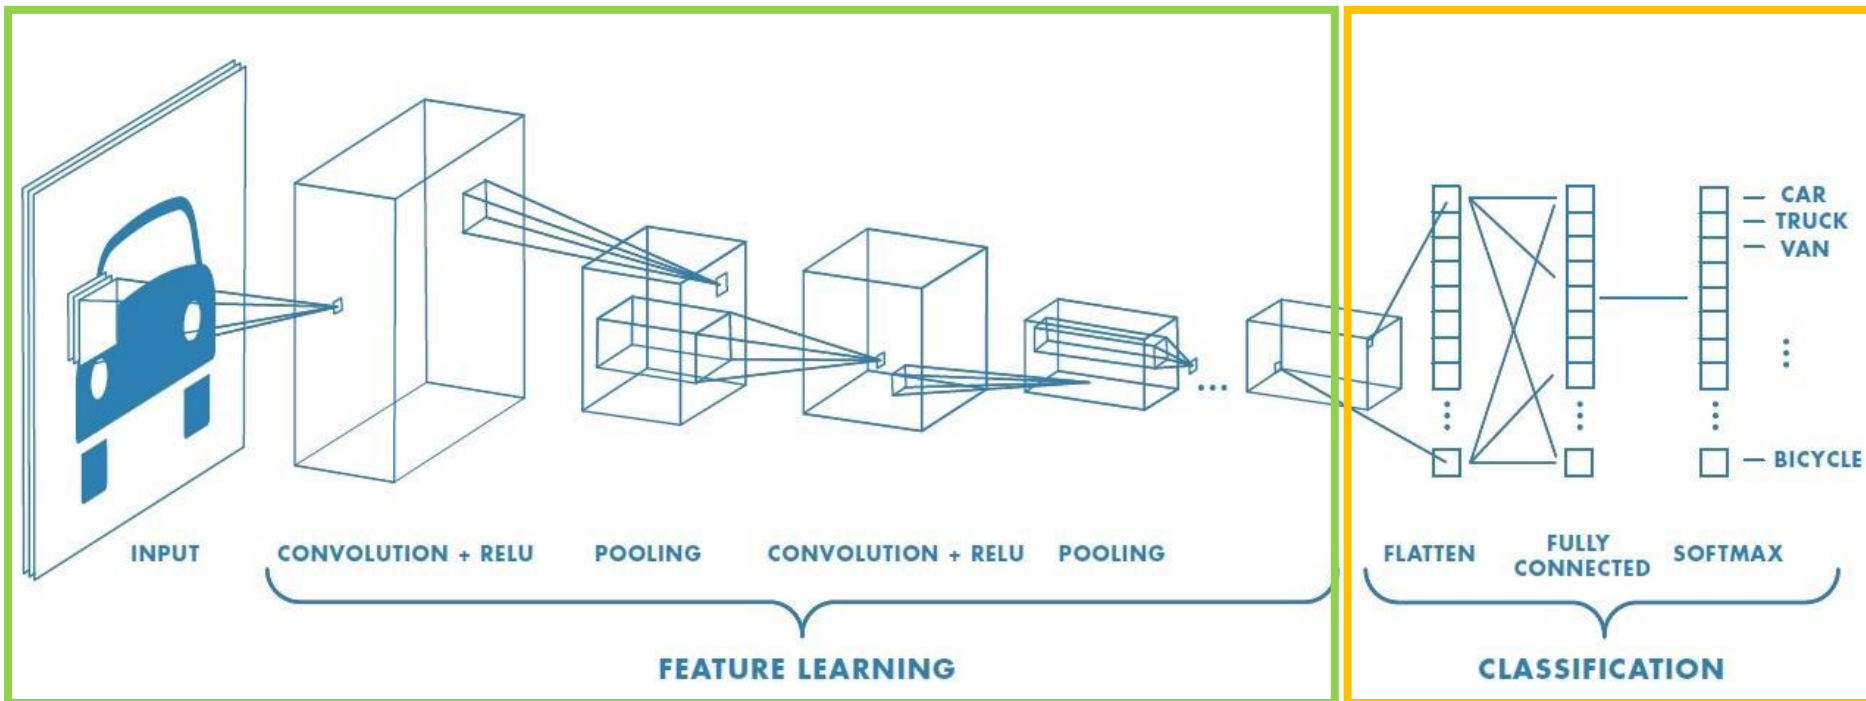
outline



output image

# Convolutional neural network

# Convolutional neural network

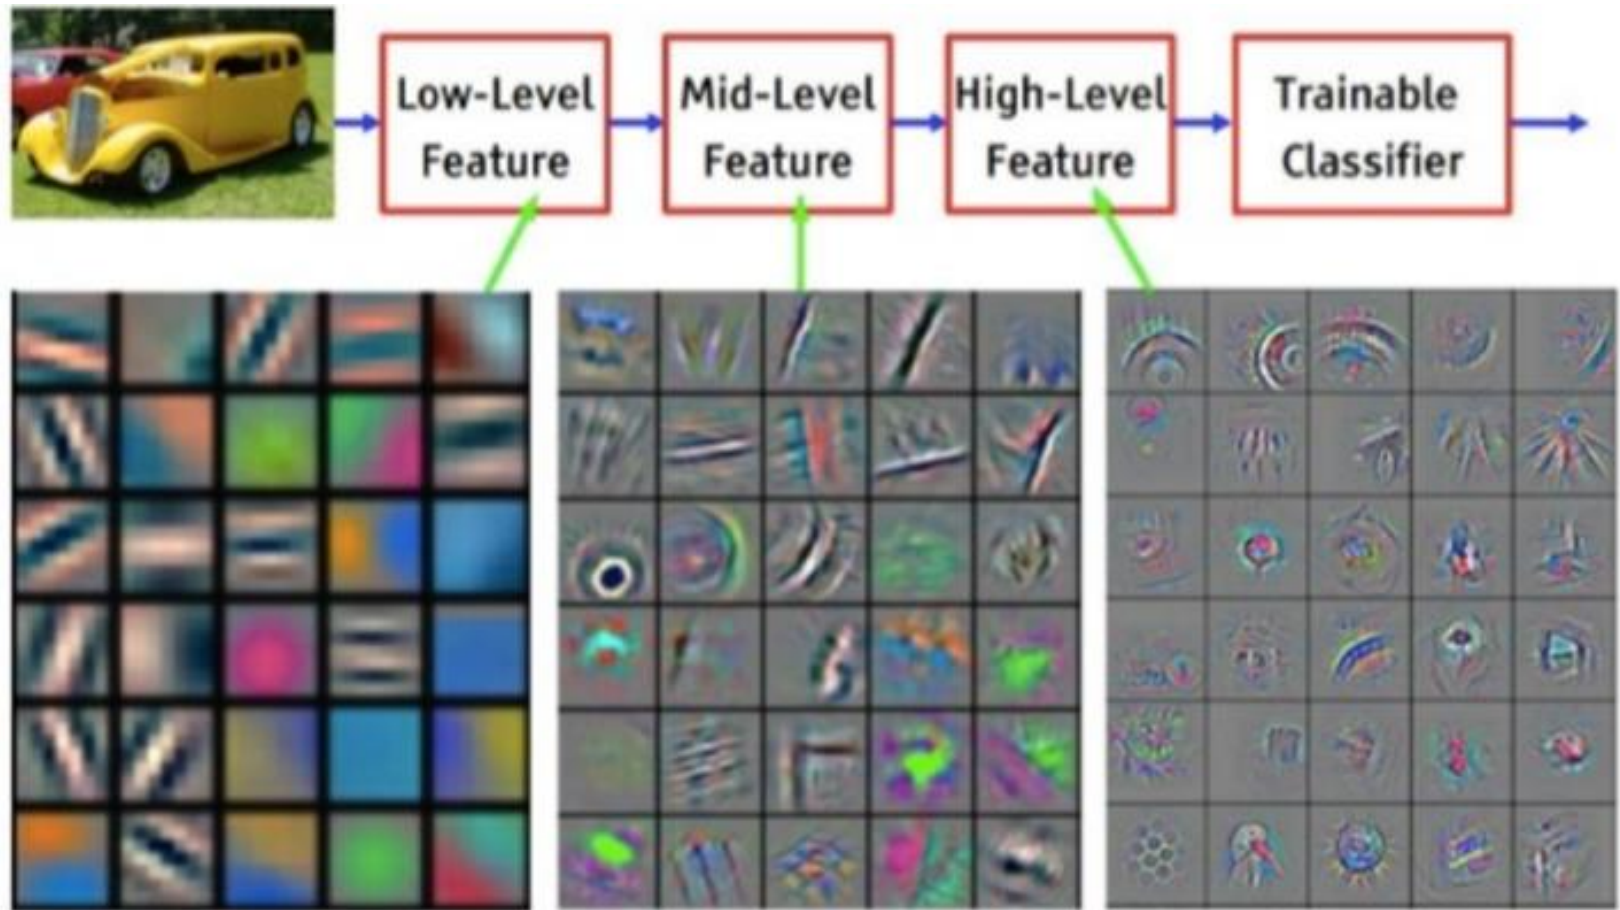


Source: towardsdatascience.com by Saha, S.

- Data-driven, automatic search for optimal filters
  - Similar objective as radiomics, but no hand-crafting
- CNN layers extract features + fully-connect layers performing prediction



# Capability of deep CNN



Source: Zeiler and Fergus (2013)

- Early layers produce simple geometric features
- Deeper layers produce complex shapes and patterns

# Backpropagation through CNN

- Input: 1D signal  $[x_1, x_2, \dots, x_d]$
- Convolution: 1x2 filter  $[w_1, w_2]$
- Output of CNN:
  - $y_1 = x_1 w_1 + x_2 w_2$
  - $y_2 = x_2 w_1 + x_3 w_2$
  - ...
  - $y_{d-1} = x_{d-1} w_1 + x_d w_2$
- Output of network:
  - $\hat{y} = \text{sigmoid}(y_1 u_1 + y_2 u_2 + \dots + y_{d-1} u_{d-1})$
- Loss:  $L(\hat{y}, y)$
- We can use chain rule to compute  $\frac{\delta L}{\delta w_1}$  and  $\frac{\delta L}{\delta w_2}$

# Convolutional layer in Tensorflow

tf.keras.layers.Conv2D

✓ See Stable

See Nightly



TensorFlow 1 version



View source on GitHub

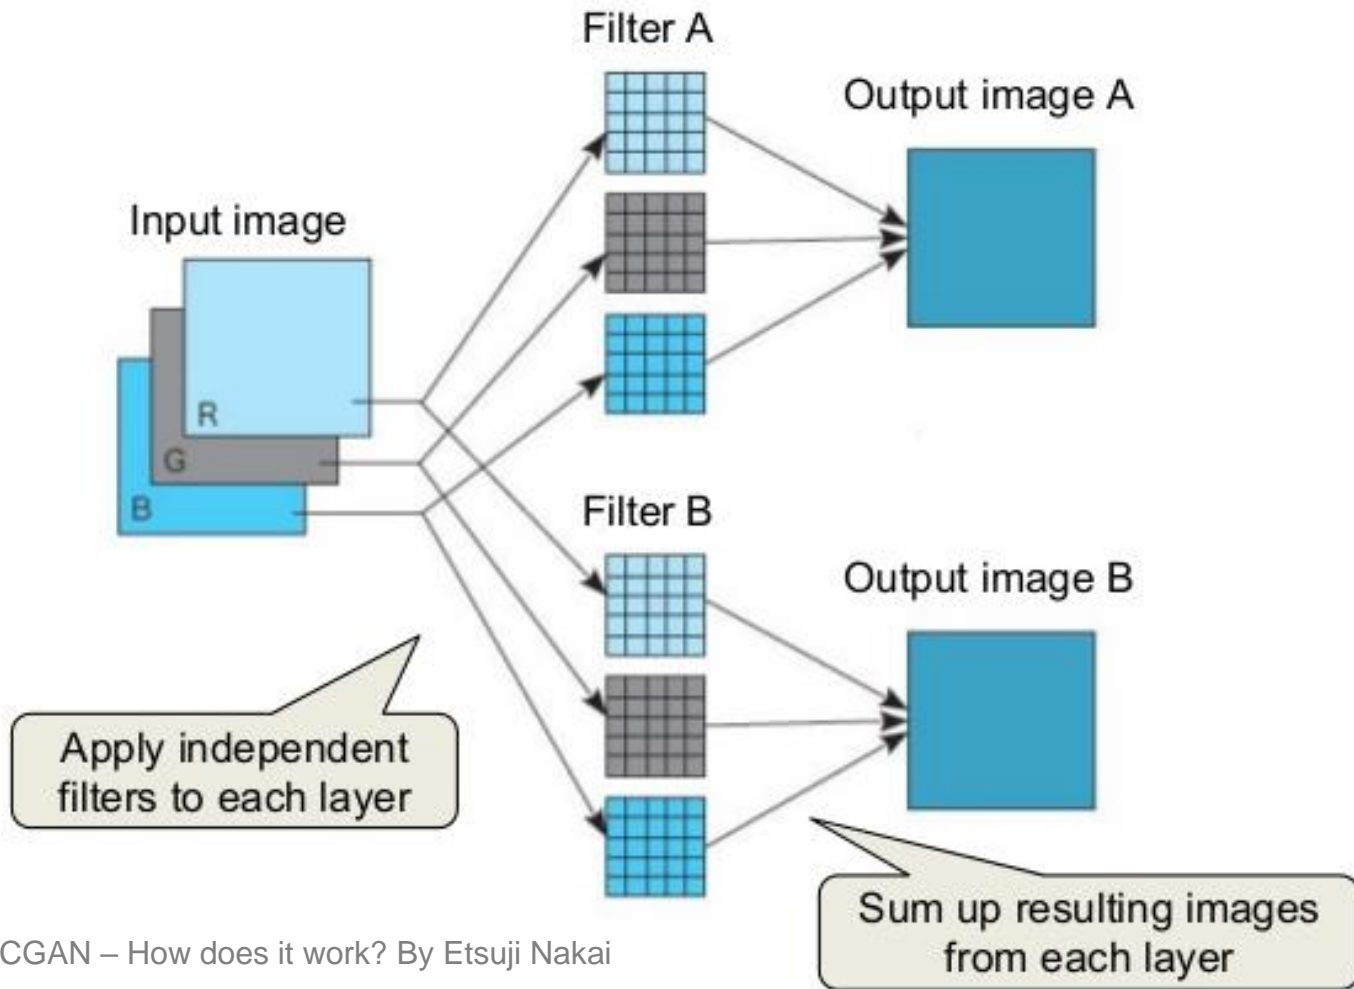
2D convolution layer (e.g. spatial convolution over images).

Inherits From: [Layer](#), [Module](#)

 [View aliases](#)

```
tf.keras.layers.Conv2D(  
    filters, kernel_size, strides=(1, 1), padding='valid',  
    data_format=None, dilation_rate=(1, 1), groups=1, activation=None,  
    use_bias=True, kernel_initializer='glorot_uniform',  
    bias_initializer='zeros', kernel_regularizer=None,  
    bias_regularizer=None, activity_regularizer=None, kernel_constraint=None,  
    bias_constraint=None, **kwargs  
)
```

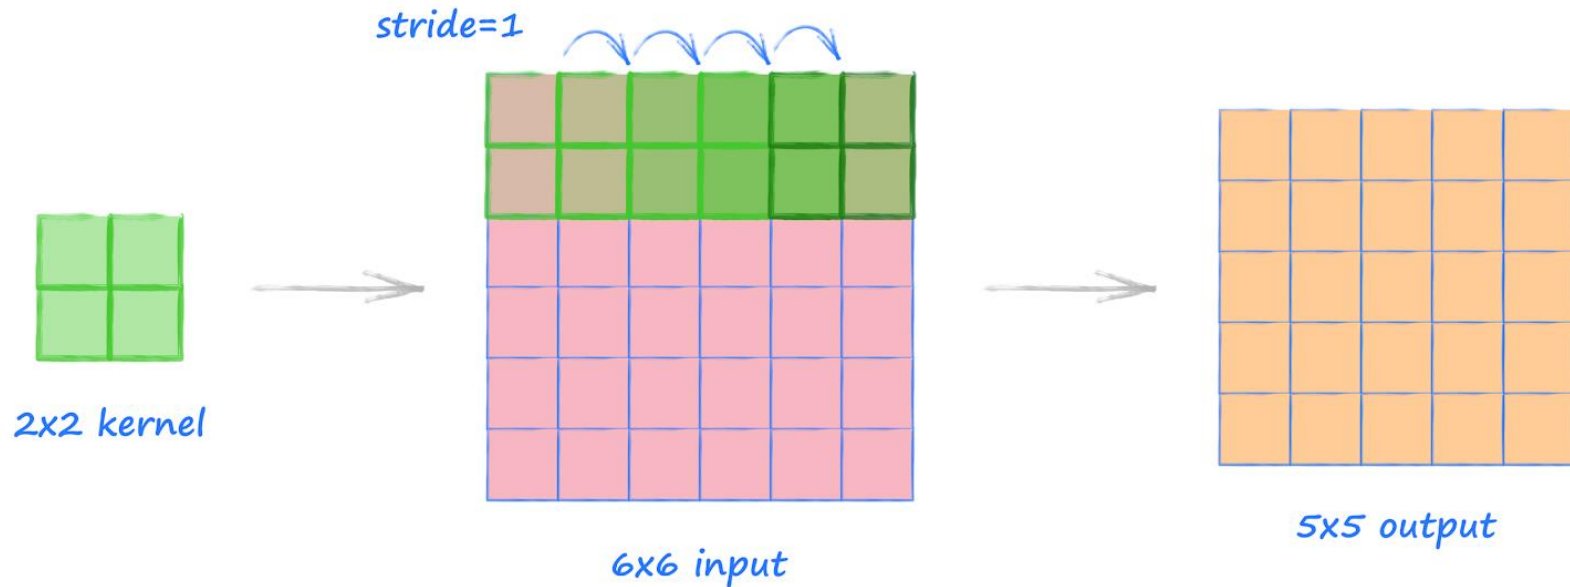
# Number of filters



Source: DCGAN – How does it work? By Etsuji Nakai

- More filters generates more geometric features

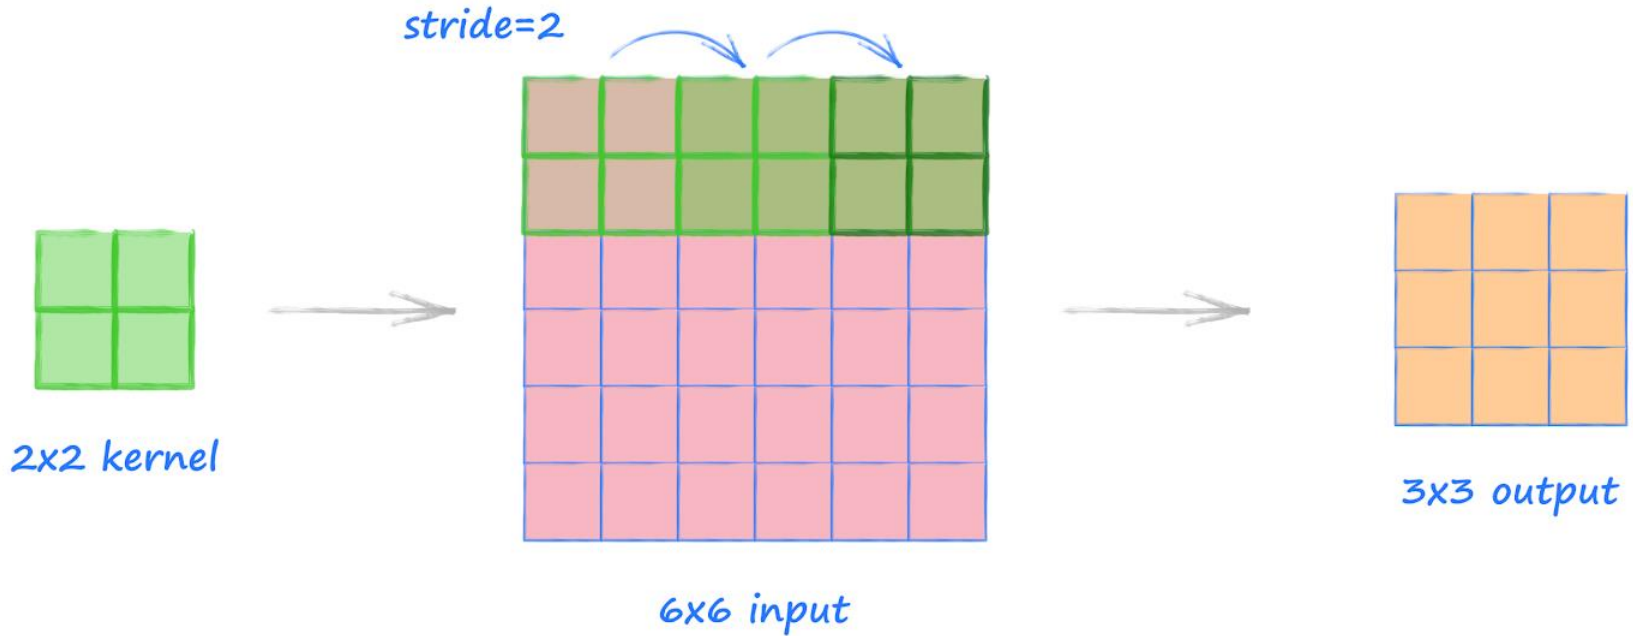
# Kernel size and stride



Source: [makeyourownneuralnetwork.blogspot.com](http://makeyourownneuralnetwork.blogspot.com)

- The size of each filter or kernel is  $k \times k$ 
  - Larger filter incorporates broader contextual data
- Stride = frequency at which to apply a filter

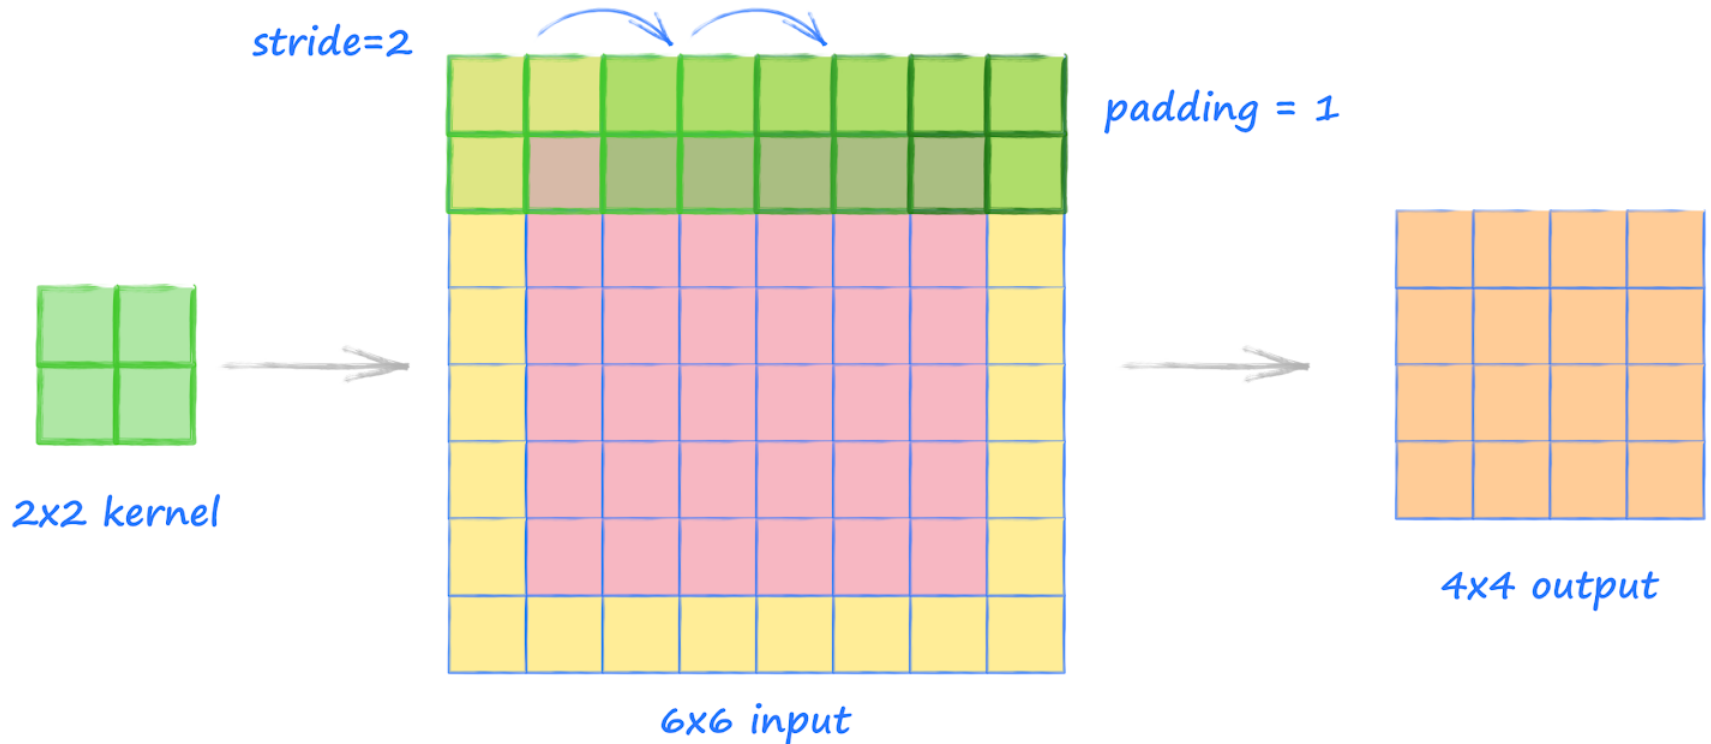
# Stride



Source: [makeyourownneuralnetwork.blogspot.com](http://makeyourownneuralnetwork.blogspot.com)

- Larger stride further reduces the output size in exchange for some loss in resolution

# Padding



Source: [makeyourownneuralnetwork.blogspot.com](http://makeyourownneuralnetwork.blogspot.com)

- **Padding = filling in zeros to the exterior of image**
  - padding = "valid" → no padding
  - padding = "same" → output size = input size

# Number of parameters in CNN

## ■ Setting A

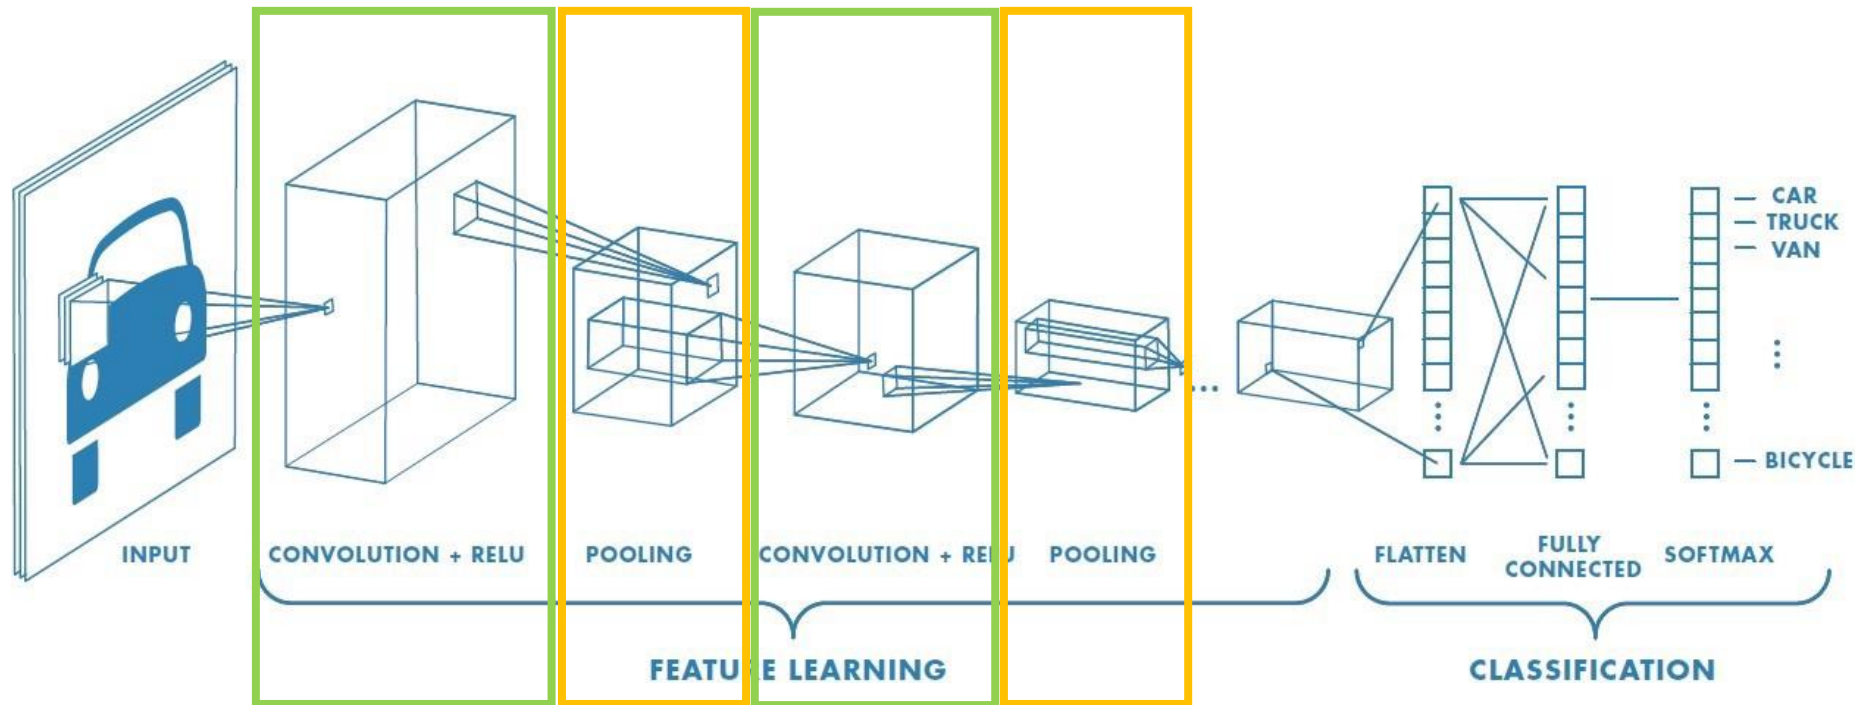
- `tf.keras.layers.Conv2D(filters = 4, kernel_size = (9, 9), strides = (1, 2), padding = 'valid')`
- Input image size = 128 x 128

## ■ Setting B

- `tf.keras.layers.Conv2D(filters = 32, kernel_size = (3, 3), strides = (1, 2), padding = 'valid')`
- `tf.keras.layers.Conv2D(filters = 8, kernel_size = (3, 3), strides = (1, 2), padding = 'valid')`
- Input image size = 64 x 64



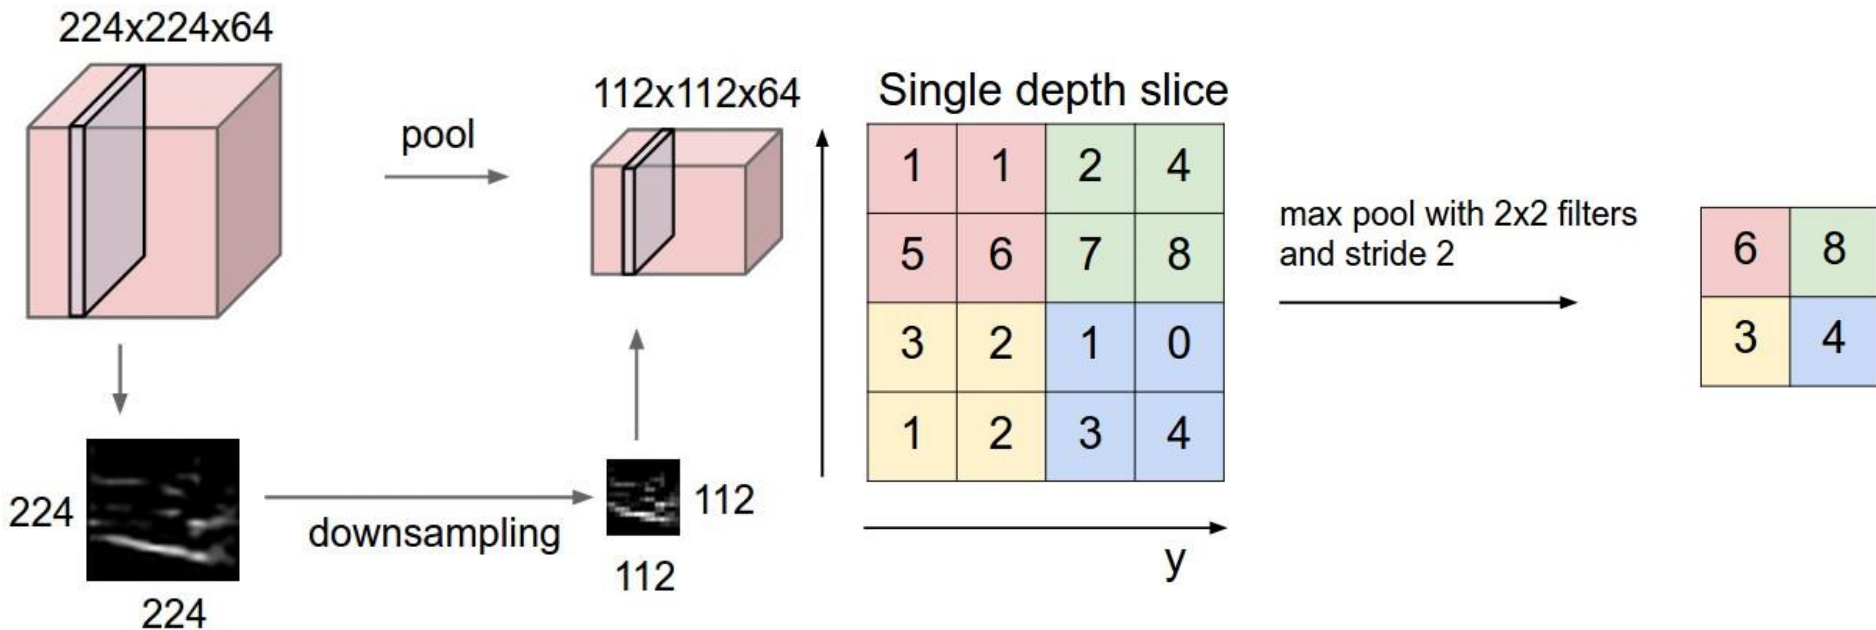
# CNN block architecture



Source: towardsdatascience.com by Saha, S.

- Convolutional layer = compute feature
- Activation layer = apply non-linear transformation
- Pooling layer = reduce dimension

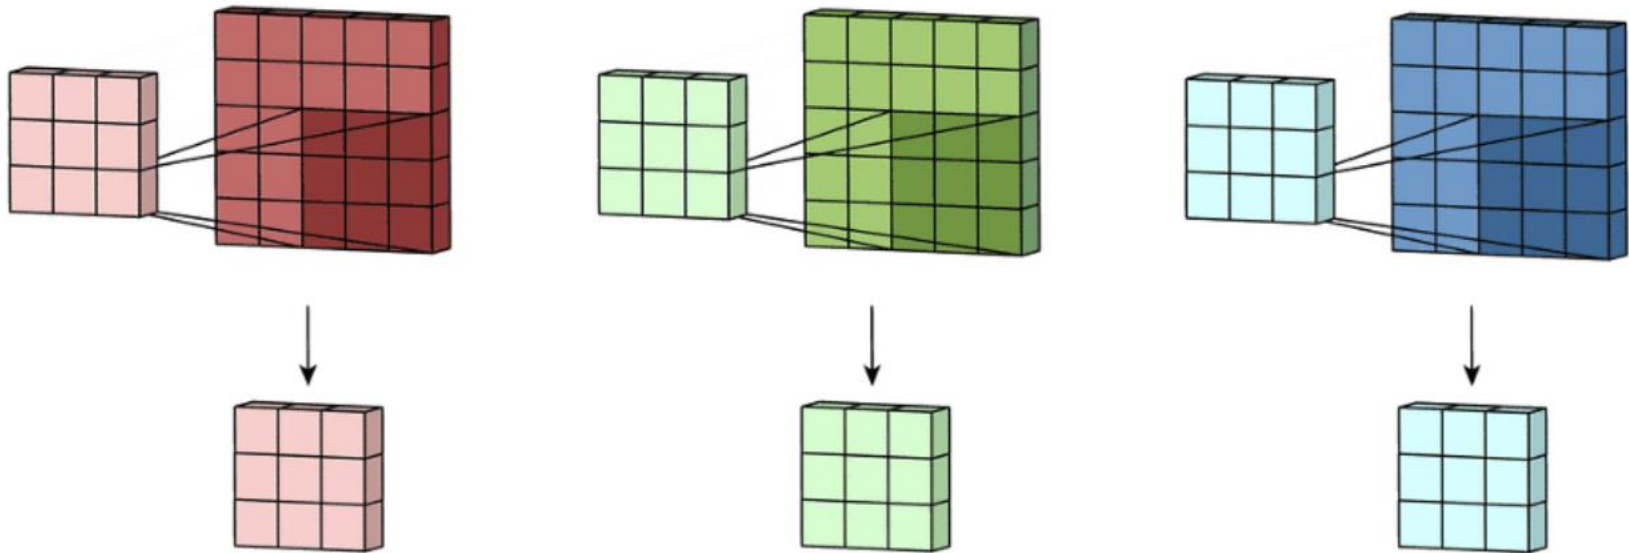
# Pooling



Source: cs231n.github.io

- Like stride, pooling layer also reduce output size
  - Max pooling
  - Average pooling

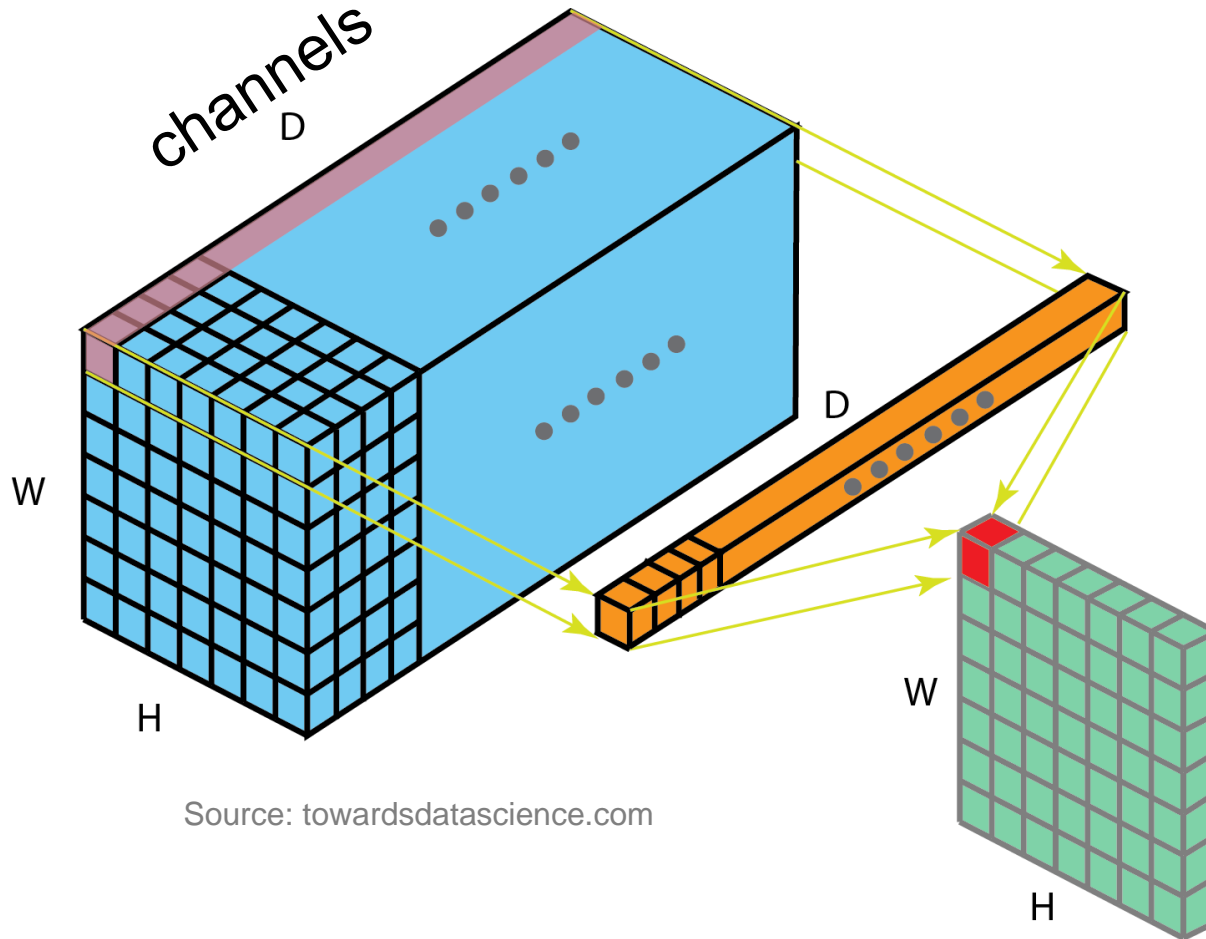
# Channel



Source: [towardsdatascience.com](https://towardsdatascience.com)

- Image can have multiple channels (like RGB)
- Each filter is applied to each channel separately
- Output of each filter is also its own channel

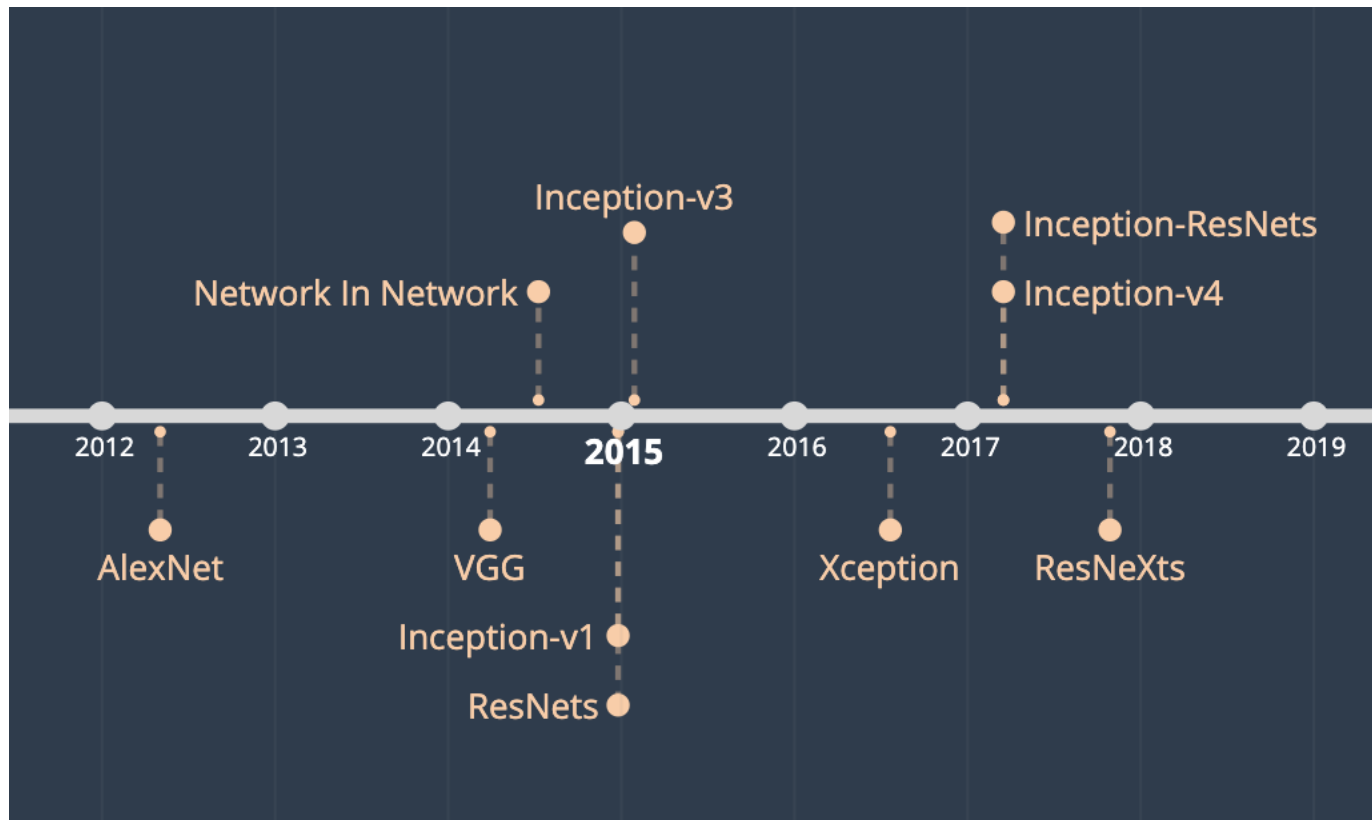
# Merging channels



- Combine values across channels for each pixel
  - Like a  $D$ -to-1 fully connected layer on channels
  - Or a  $1 \times 1$  convolution on image

# Key CNN architectures

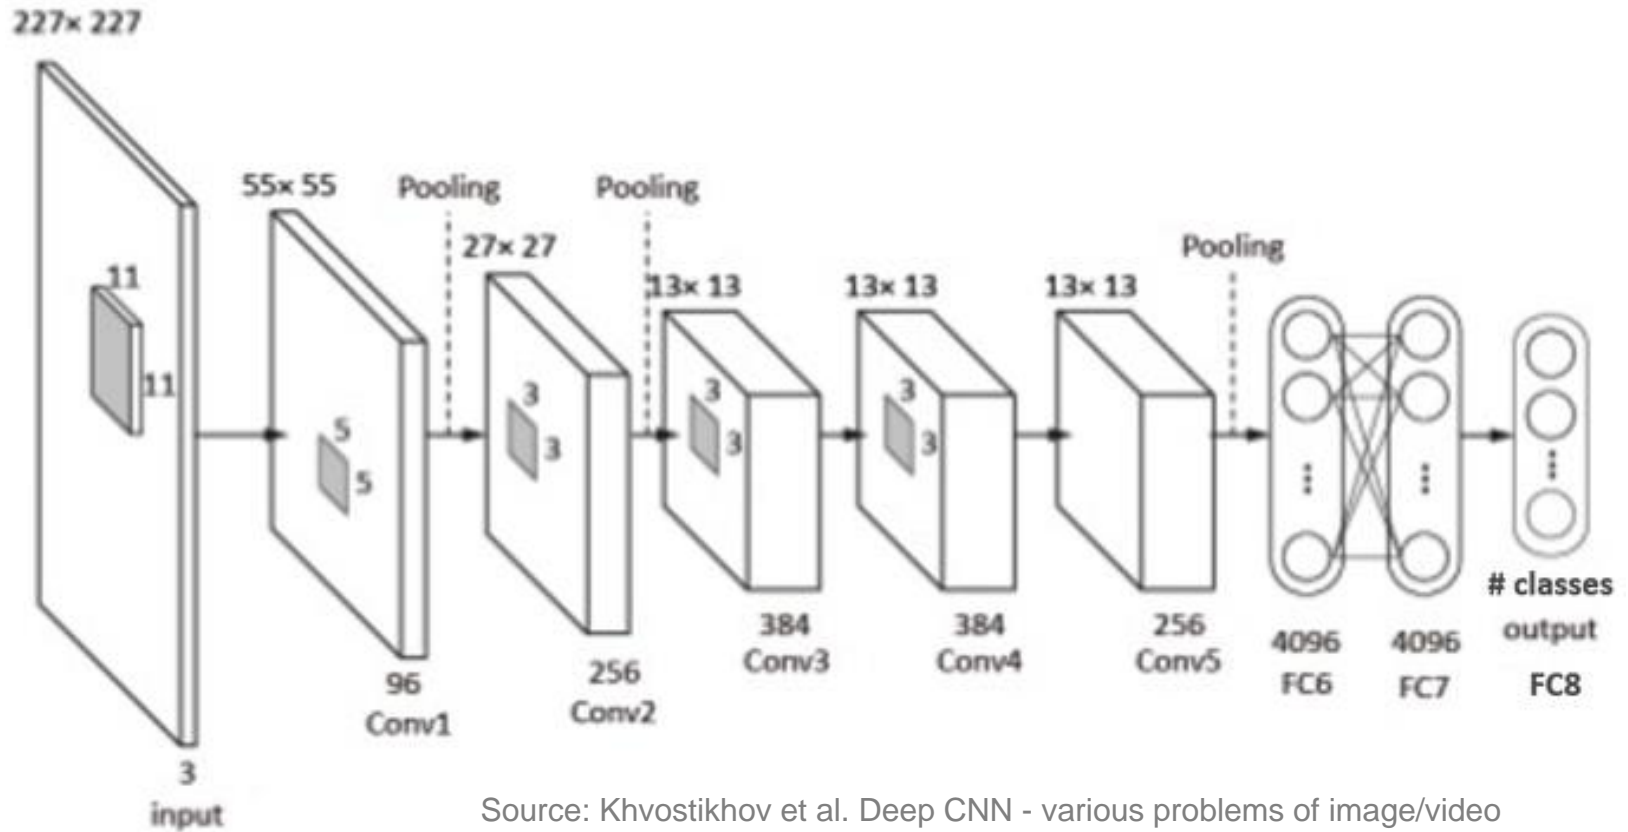
# Timeline of CNN (for classification)



Source: [towardsdatascience.com](https://towardsdatascience.com)

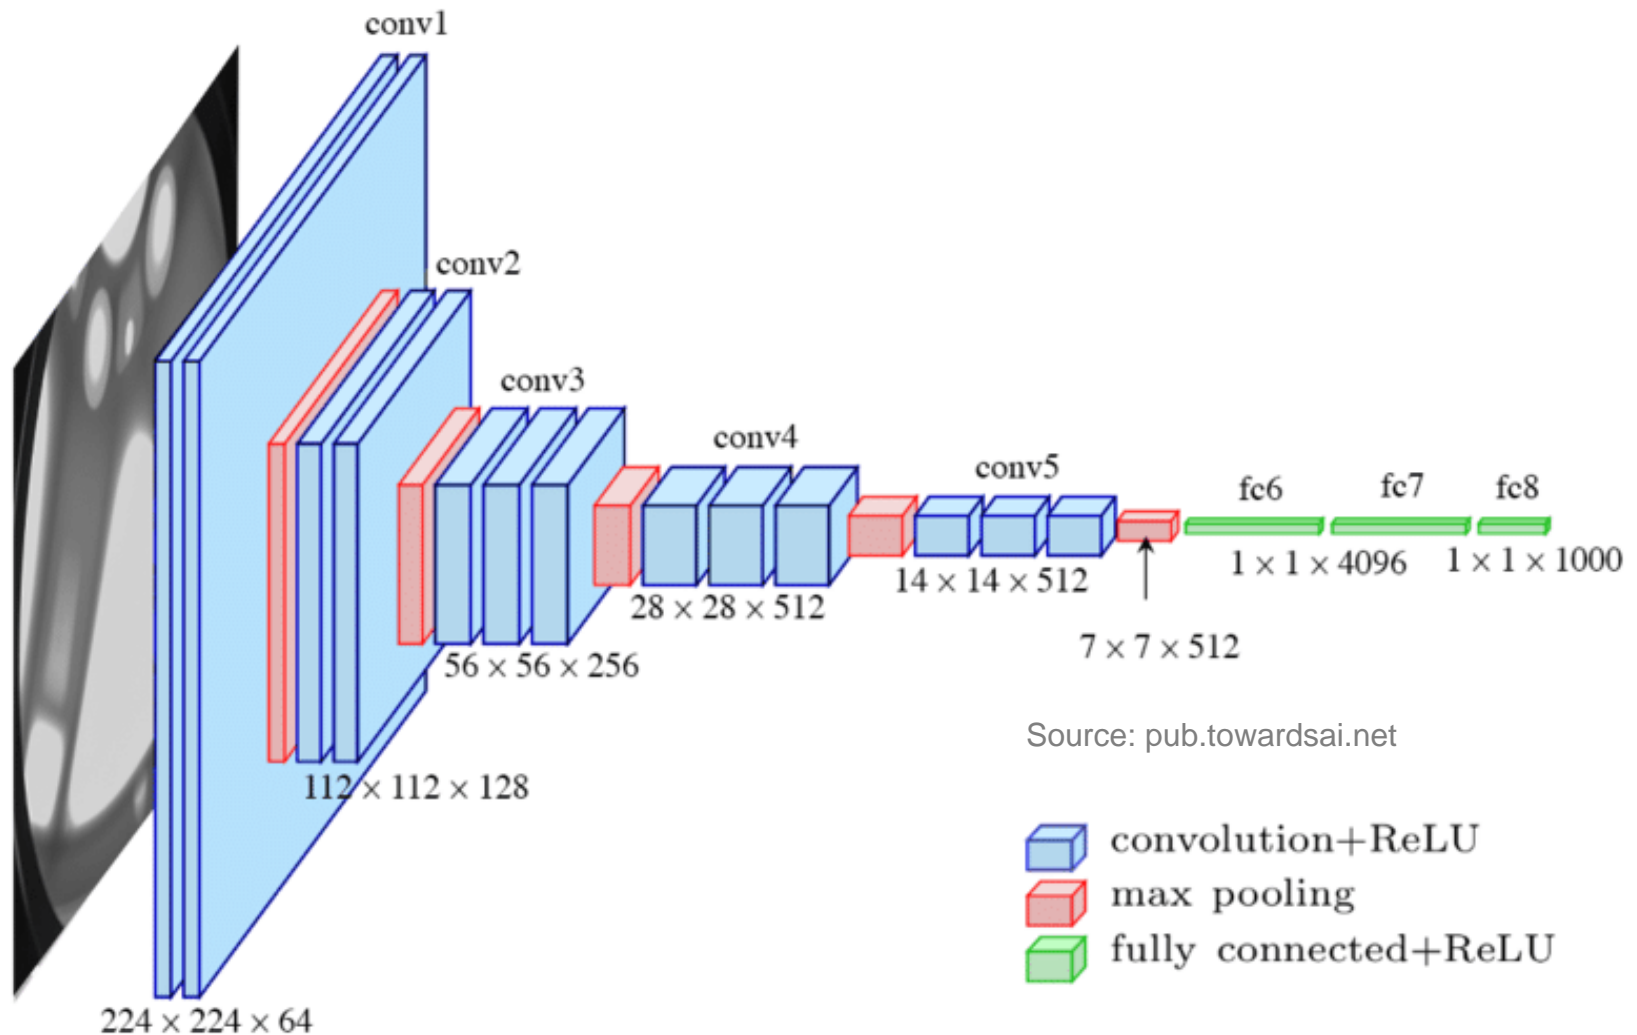
- AlexNet was the winner of 2012 ImageNet challenge
  - 5 convolution + 3 fully connected layers
  - Better than 2<sup>nd</sup> place by more than 10 percentage points

# AlexNet



Source: Khvostikhov et al. Deep CNN - various problems of image/video classification Alzheimer disease studies (2018)

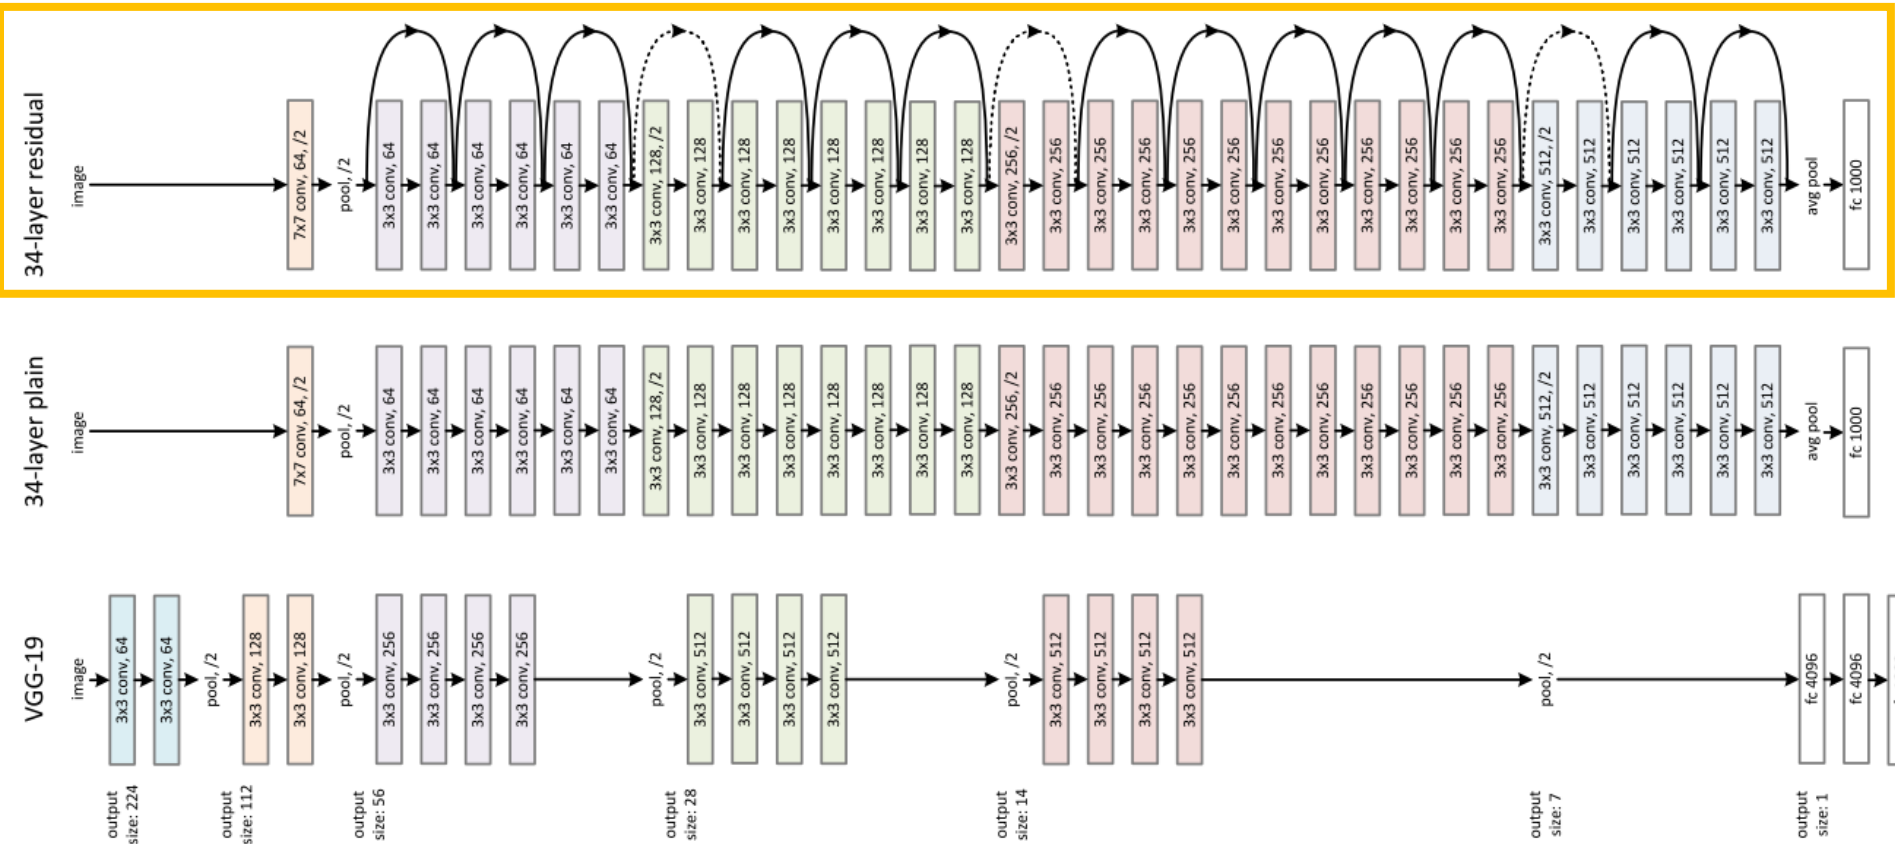
# VGG-16



- Decrease image size, increase number of channels



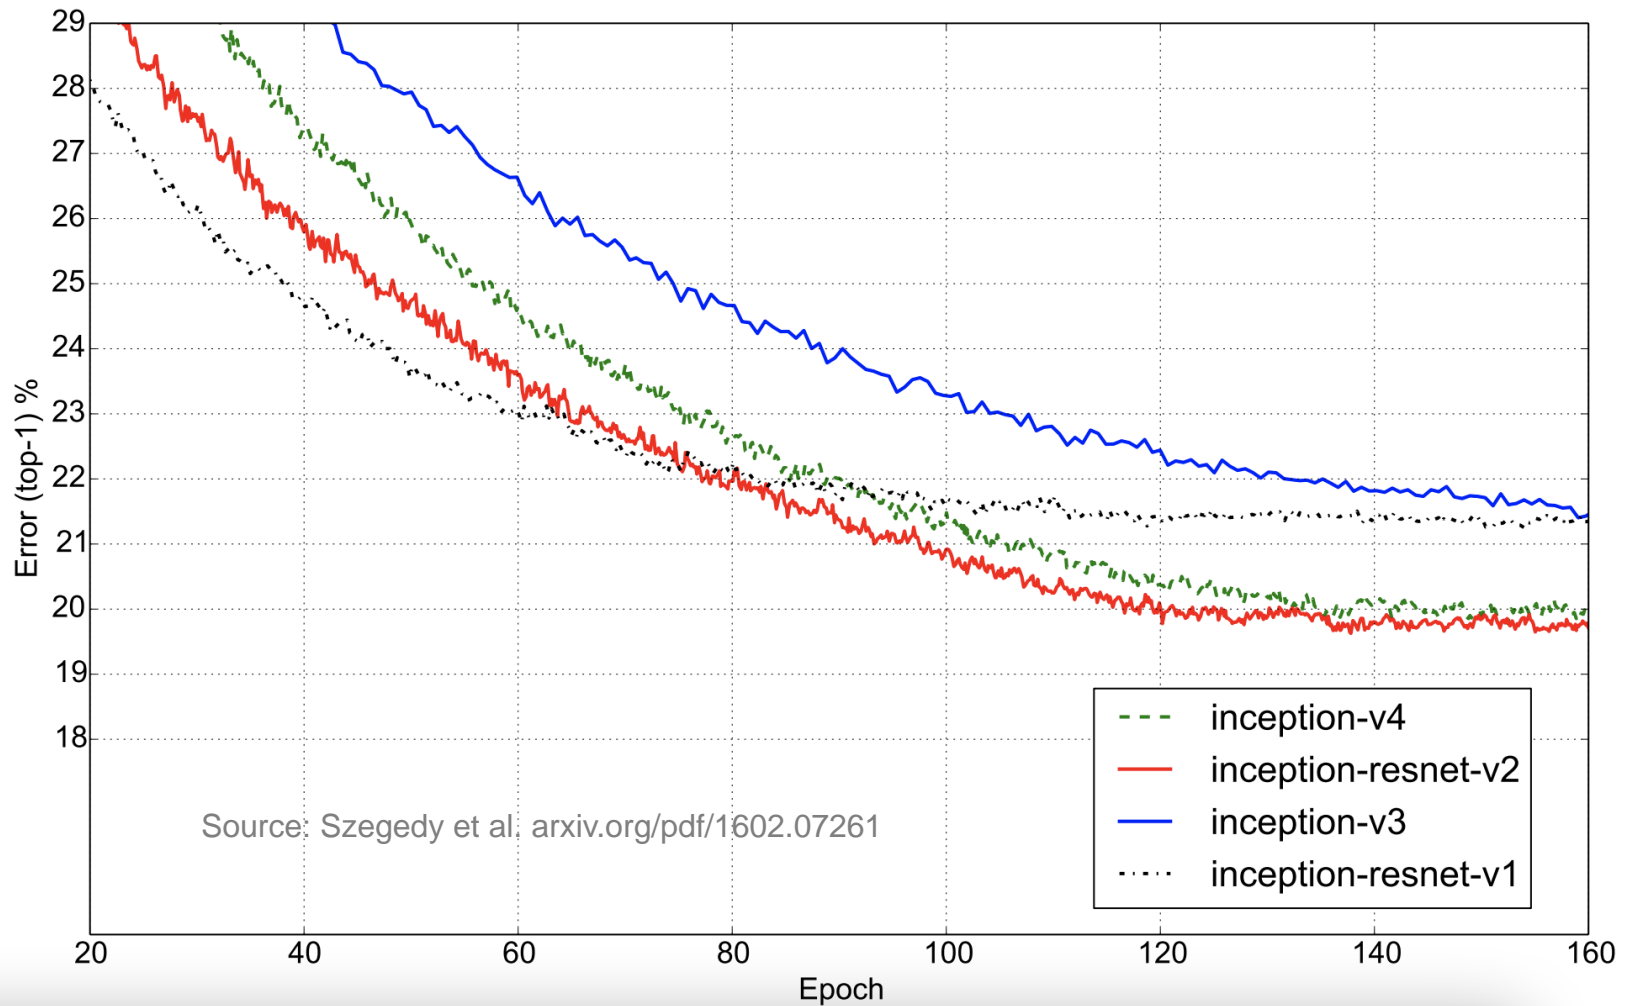
# Residual Network (ResNet)



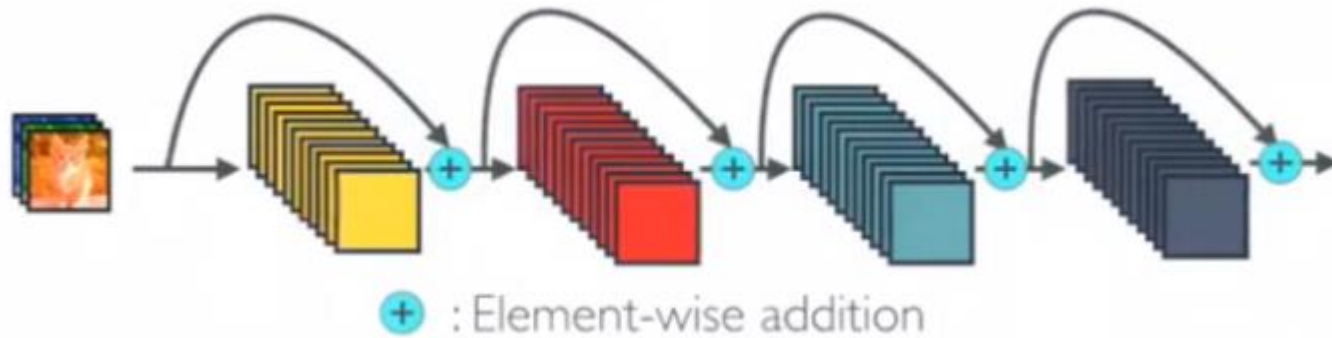
Source: medium.com

- Add bypasses over multiple convolutional layers
  - Reduce the number of multiplication terms in chain rule during backpropagation for early layers

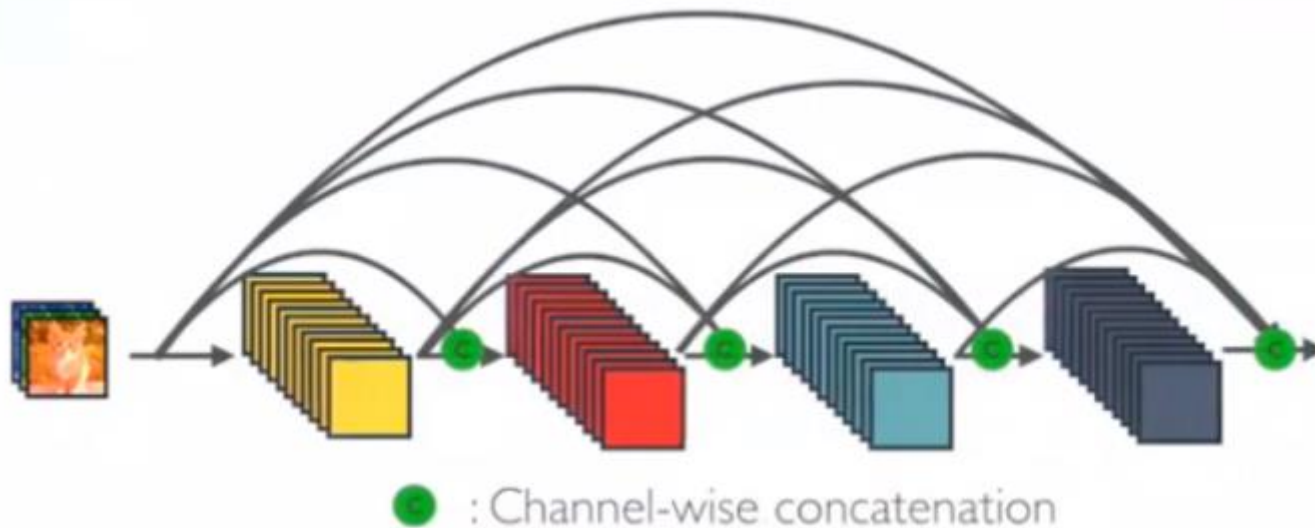
# Residual connection improves training



# DenseNet

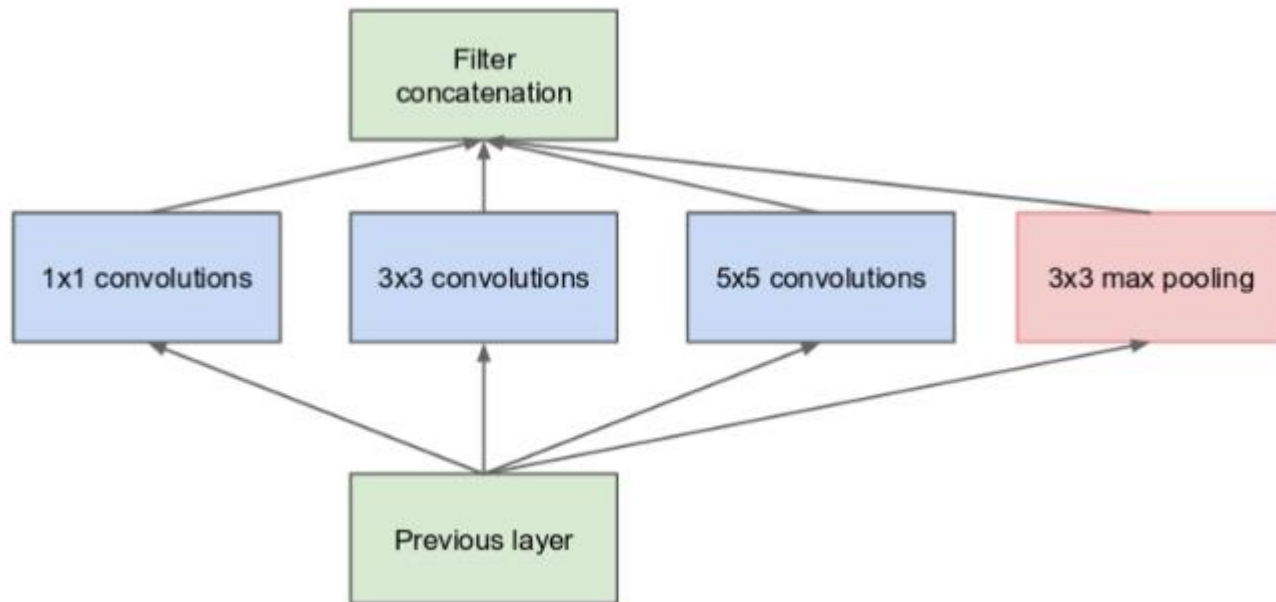


Source: towardsdatascience.com



- Feedforward bypasses

# Inception = multi-resolution CNN



From left: A dog occupying most of the image, a dog occupying a part of it, and a dog occupying very little space (Images obtained from [Unsplash](#)).

# CNN design in a nutshell



# Module: tf.keras.applications

## Modules

`densenet` module: DenseNet models for Keras.

`efficientnet` module: EfficientNet models for Keras.

`imagenet_utils` module: Utilities for ImageNet data preprocessing & prediction decoding.

`inception_resnet_v2` module: Inception-ResNet V2 model for Keras.

`inception_v3` module: Inception V3 model for Keras.

`mobilenet` module: MobileNet v1 models for Keras.

`mobilenet_v2` module: MobileNet v2 models for Keras.

`mobilenet_v3` module: MobileNet v3 models for Keras.

`nasnet` module: NASNet-A models for Keras.

`resnet` module: ResNet models for Keras.

# Any question?



Pyception video, AnacondaCon 2018