

# 3011979 Intro to Deep Learning for Medical Imaging

## L12: Neural network training & image segmentation architecture

Apr 23<sup>rd</sup>, 2021



Sira Sriswasdi, Ph.D.

Research Affairs, Faculty of Medicine  
Chulalongkorn University

# Advanced ANN training techniques

# Representation

	1	2	3	4	5	6	7	8	9
man	1	0	0	0	0	0	0	0	0
woman	0	1	0	0	0	0	0	0	0
boy	0	0	1	0	0	0	0	0	0
girl	0	0	0	1	0	0	0	0	0
prince	0	0	0	0	1	0	0	0	0
princess	0	0	0	0	0	1	0	0	0
queen	0	0	0	0	0	0	1	0	0
king	0	0	0	0	0	0	0	1	0
monarch	0	0	0	0	0	0	0	0	1

Image from hackermoon.com


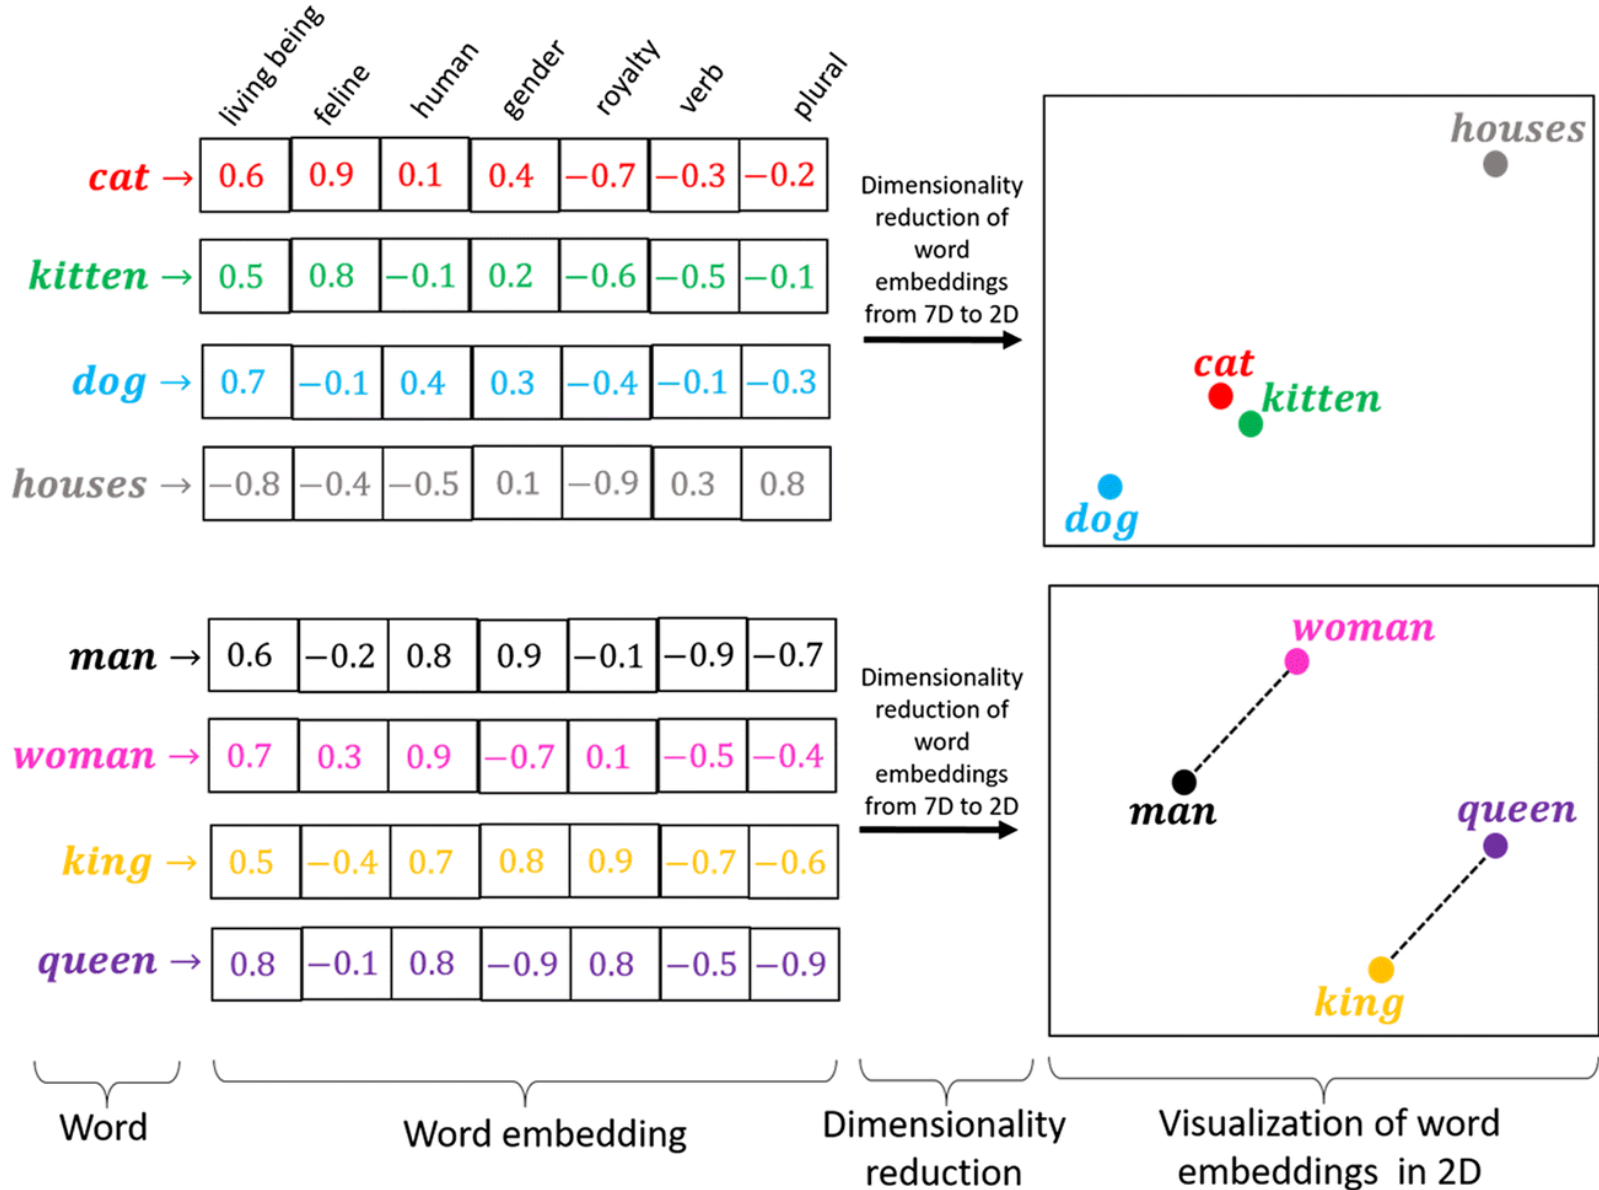


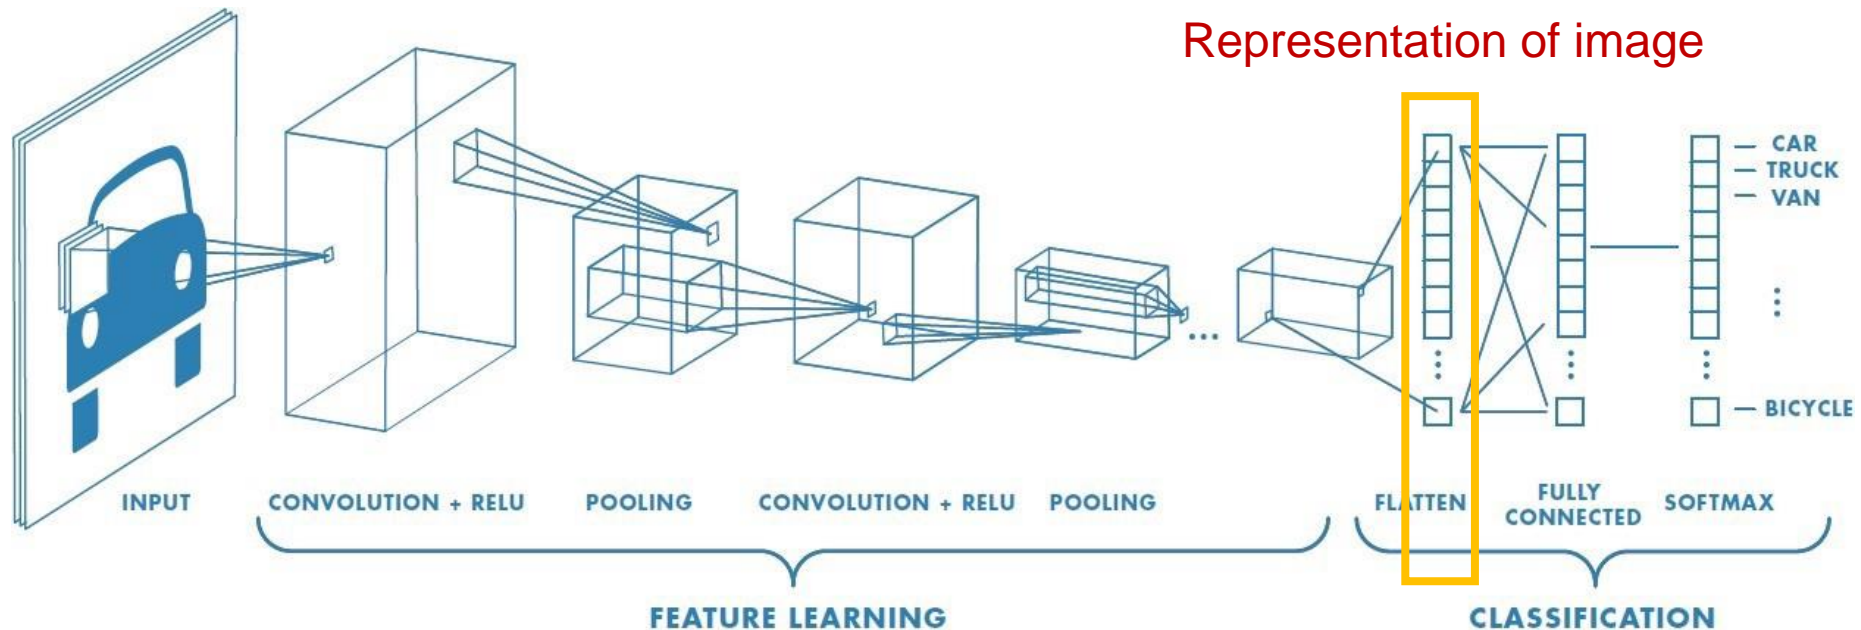
Image from naushardsblog.wordpress.com

- Representation = presentation of information contained in the data
- Simple representation is not useful for learning

# Good representation



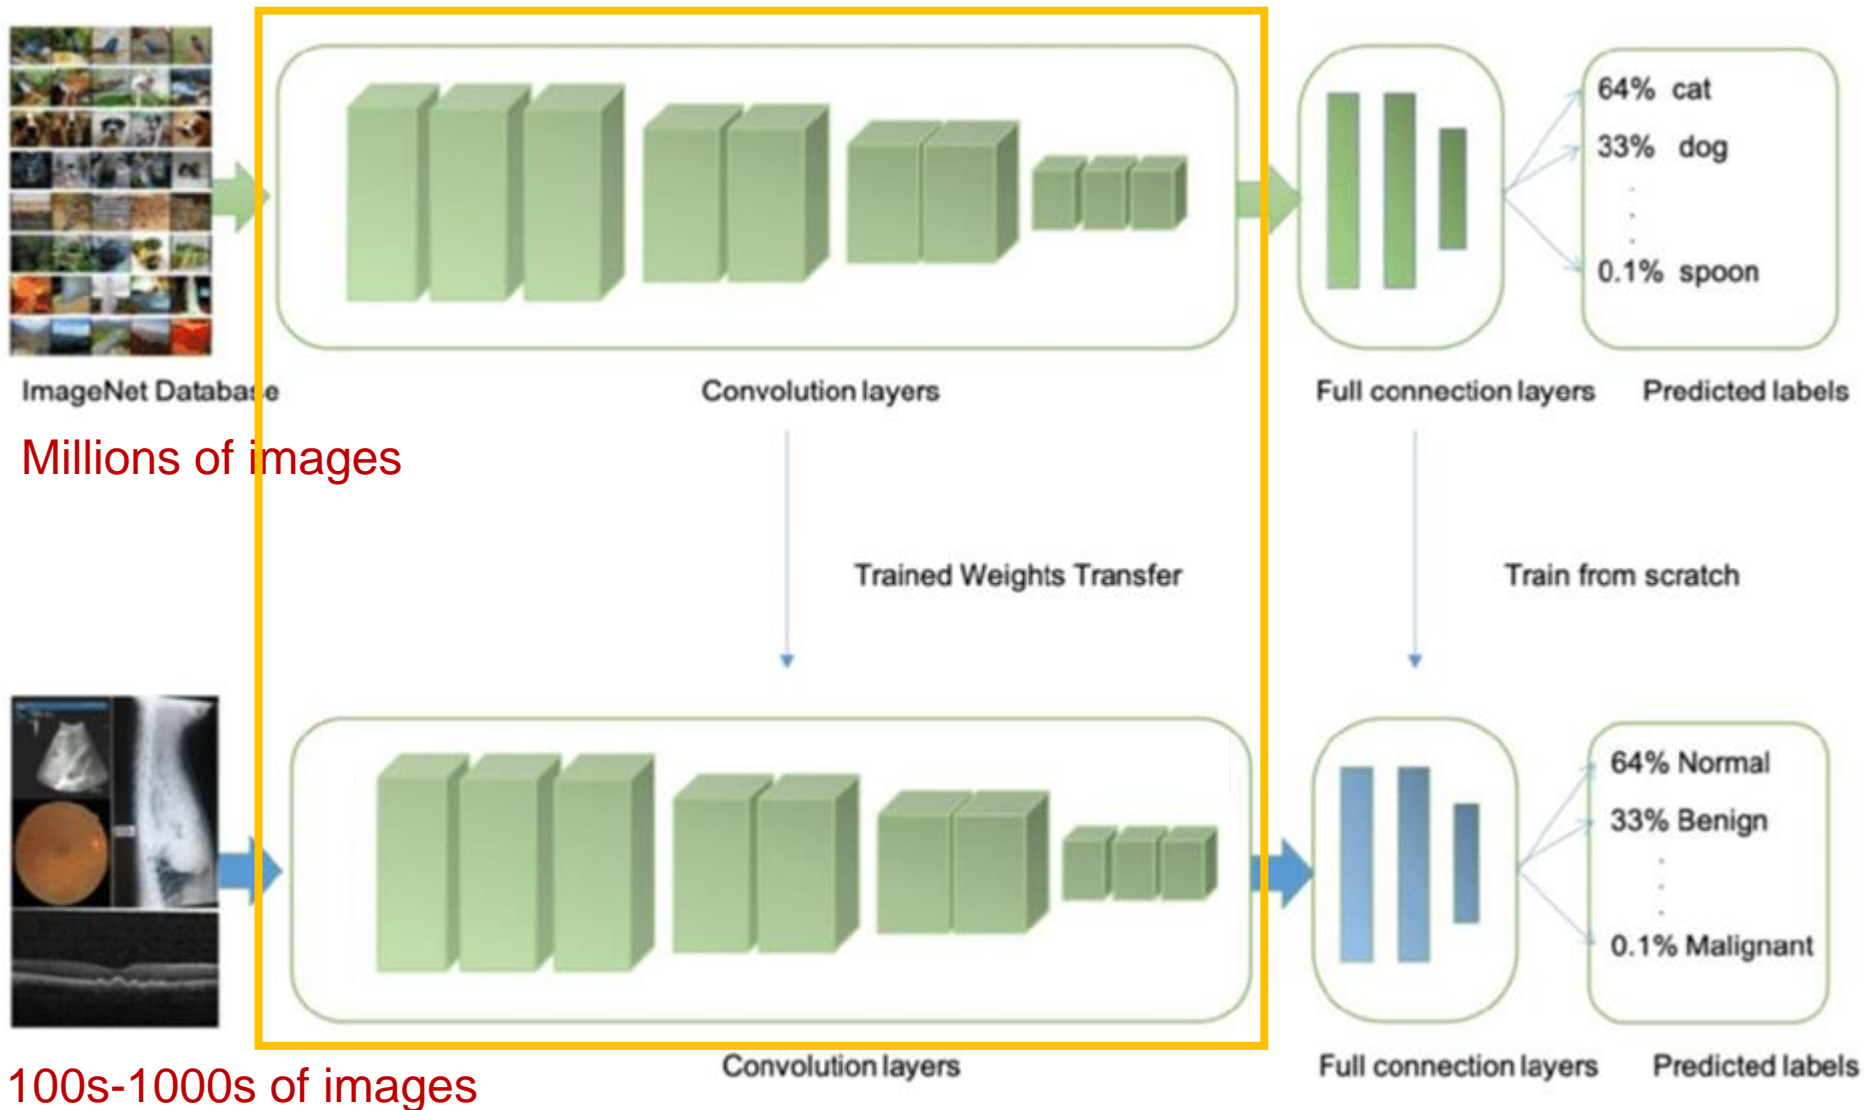
# Representation learning



Source: towardsdatascience.com by Saha, S.

- Representation learning = machine learning for obtaining a good representation of the data
  - High performance  $\leftrightarrow$  Good representation
- Convolution layers in a CNN

# Transfer learning



# Tensorflow pre-defined models

TensorFlow > API > TensorFlow Core v2.4.1 > Python

## Module: tf.keras.applications

`densenet` module: DenseNet models for Keras.

`efficientnet` module: EfficientNet models for Keras.

`imagenet_utils` module: Utilities for ImageNet data preprocessing & prediction decoding.

`inception_resnet_v2` module: Inception-ResNet V2 model for Keras.

`inception_v3` module: Inception V3 model for Keras.

`mobilenet` module: MobileNet v1 models for Keras.

`mobilenet_v2` module: MobileNet v2 models for Keras.

`mobilenet_v3` module: MobileNet v3 models for Keras.

`nasnet` module: NASNet-A models for Keras.

`resnet` module: ResNet models for Keras.

# Using pre-trained models

tf.keras.applications.DenseNet121

```
tf.keras.applications.DenseNet121(  
    include_top=True, weights='imagenet', input_tensor=None,  
    input_shape=None, pooling=None, classes=1000  
)
```

- **include\_top**
  - Whether to include the fully connected layers
  - Set to **False** if to train fully connected layers from scratch
    - Remember to add new fully connected layers
- **weights**
  - Path to pretrained model weights
  - “imagenet” is provided by default



# More pre-trained models

## tf.keras.applications.ResNet50

```
tf.keras.applications.ResNet50(  
    include_top=True, weights='imagenet', input_tensor=None,  
    input_shape=None, pooling=None, classes=1000, **kwargs  
)
```

## tf.keras.applications.InceptionResNetV2

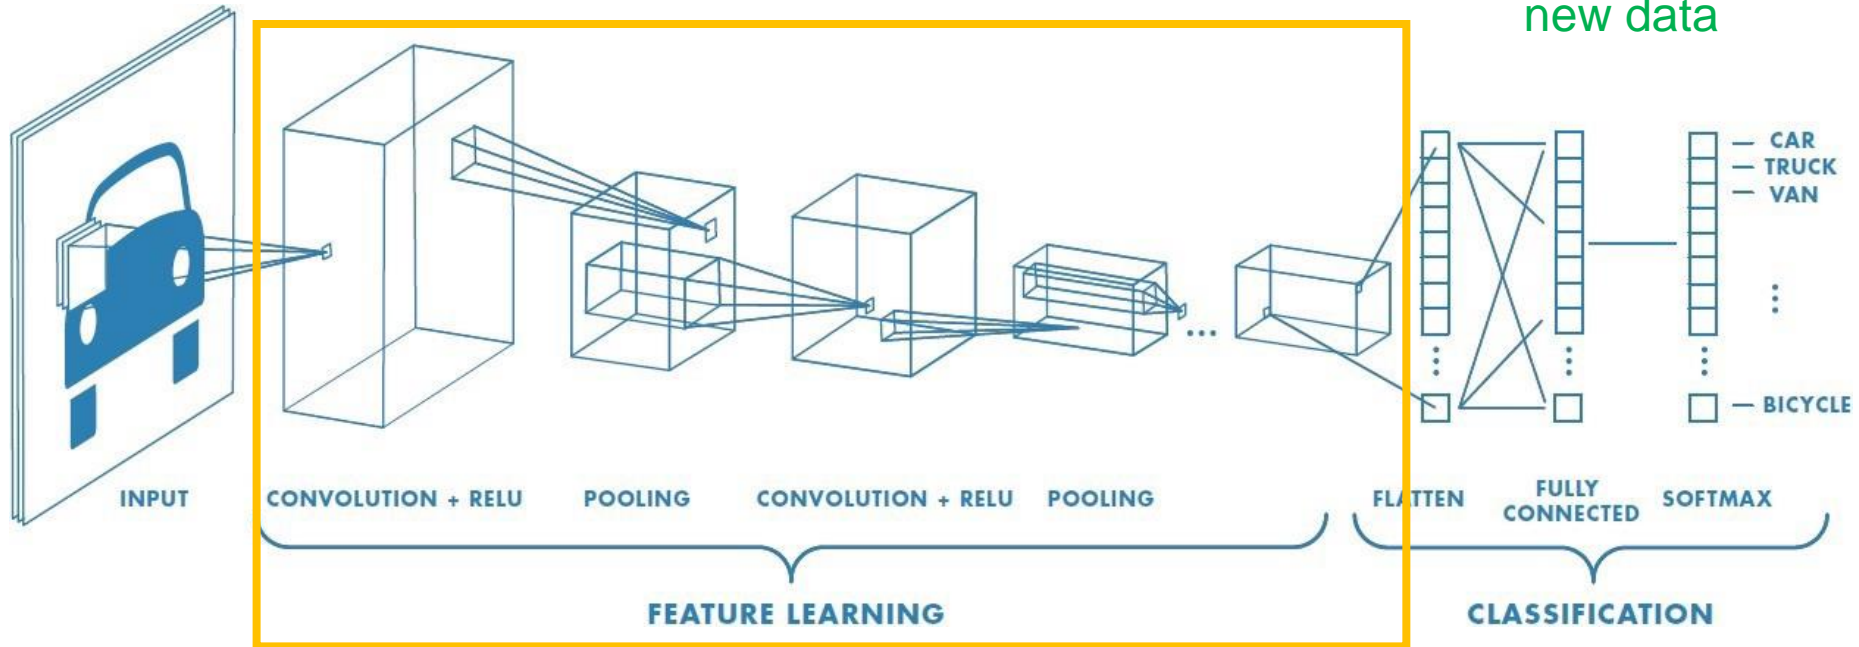
```
tf.keras.applications.InceptionResNetV2(  
    include_top=True, weights='imagenet', input_tensor=None,  
    input_shape=None, pooling=None, classes=1000,  
    classifier_activation='softmax', **kwargs  
)
```

# Transfer learning

New data

Load weights from ImageNet pretrained

Train using  
new data



- Fully connected part is untrained
- Convolution part is pretrained
  - Freeze its weights

# Layer freezing

```
base_model = keras.applications.Xception(  
    weights='imagenet', # Load weights pre-trained on ImageNet.  
    input_shape=(150, 150, 3),  
    include_top=False) # Do not include the ImageNet classifier at the top.
```

```
base_model.trainable = False
```

```
inputs = keras.Input(shape=(150, 150, 3))  
# We make sure that the base_model is running in inference mode here,  
# by passing `training=False`. This is important for fine-tuning, as you will  
# learn in a few paragraphs.  
x = base_model(inputs, training=False)  
# Convert features of shape `base_model.output_shape[1:]` to vectors  
x = keras.layers.GlobalAveragePooling2D()(x)  
# A Dense classifier with a single unit (binary classification)  
outputs = keras.layers.Dense(1)(x)  
model = keras.Model(inputs, outputs)
```

- [https://www.tensorflow.org/guide/keras/transfer\\_learning](https://www.tensorflow.org/guide/keras/transfer_learning)

# Fine-tuning


## Transfer learning

- Load pretrained model
- Freeze pretrained layers
- Attach new fully connected layers
- Train model using new dataset

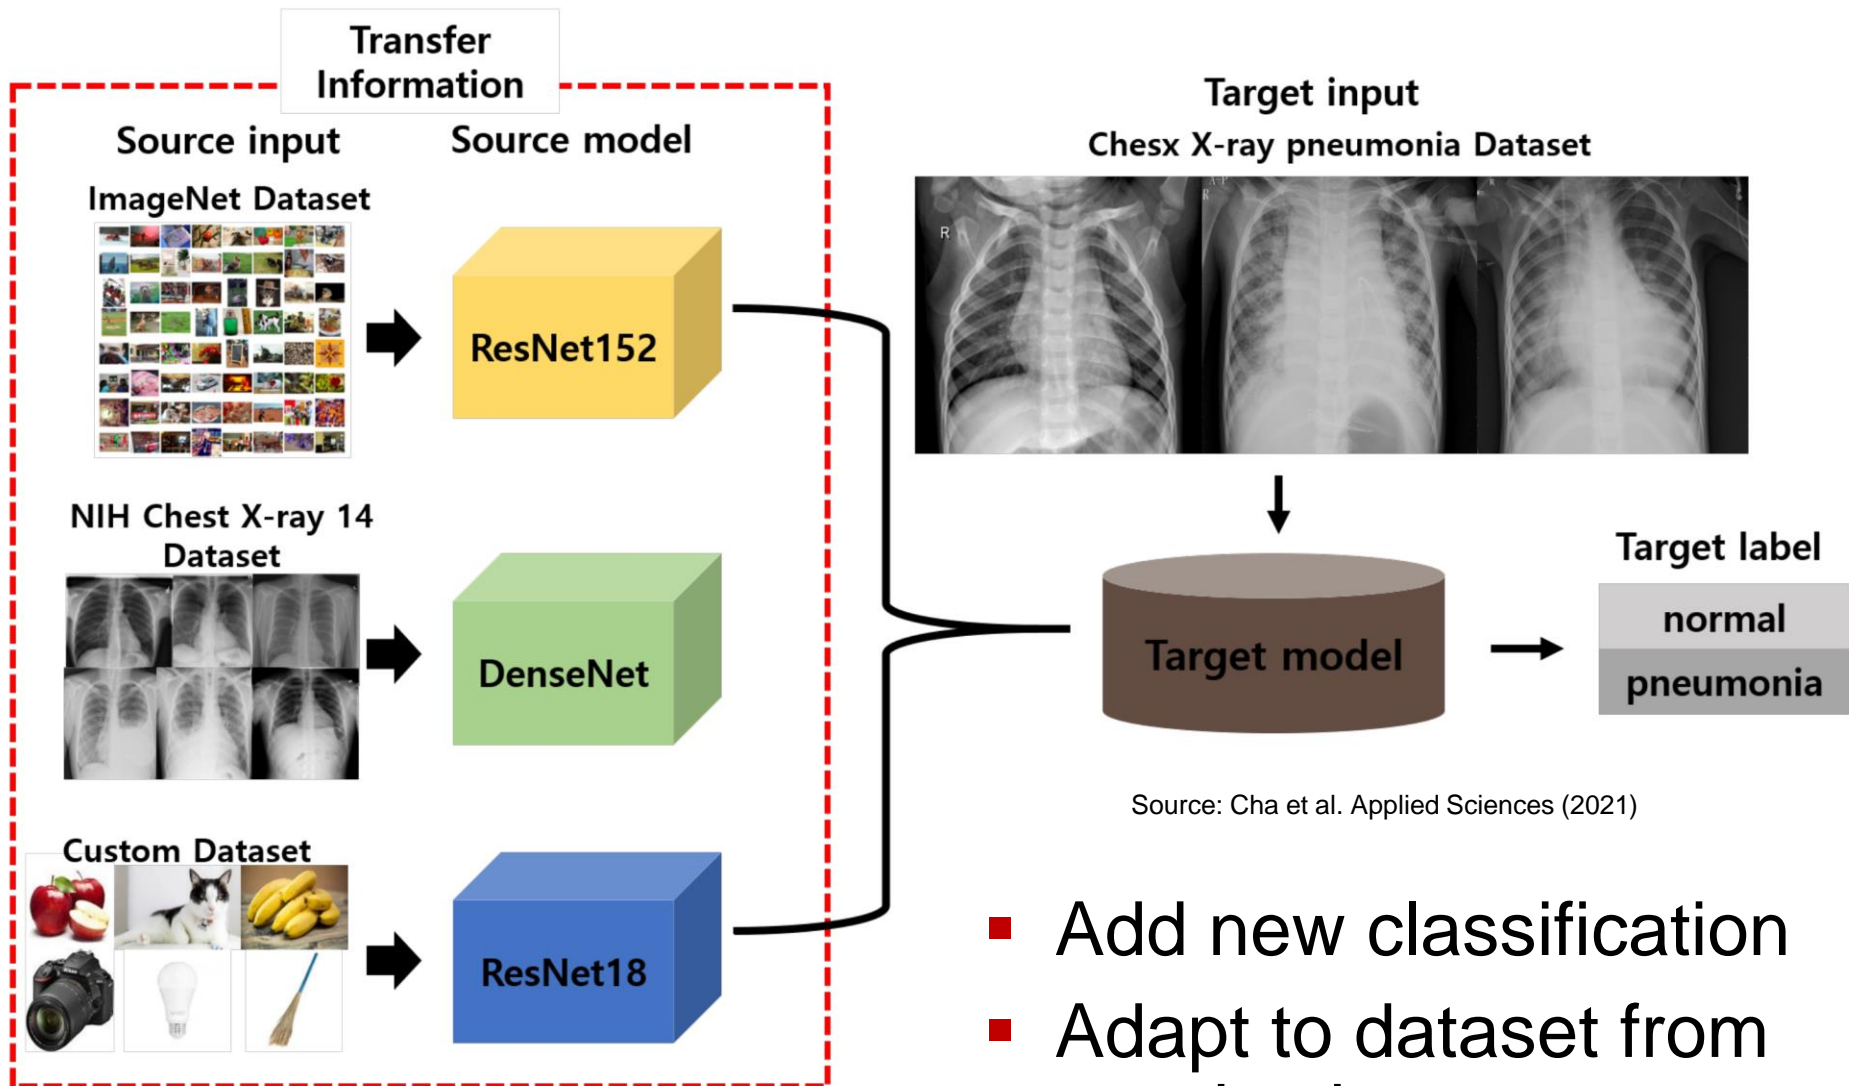
## Fine-tuning

- Unfreeze pretrained layers
- Continue to train the model using new dataset

# Impact of transfer learning




# Transfer learning application



- Add new classification
- Adapt to dataset from new institutes

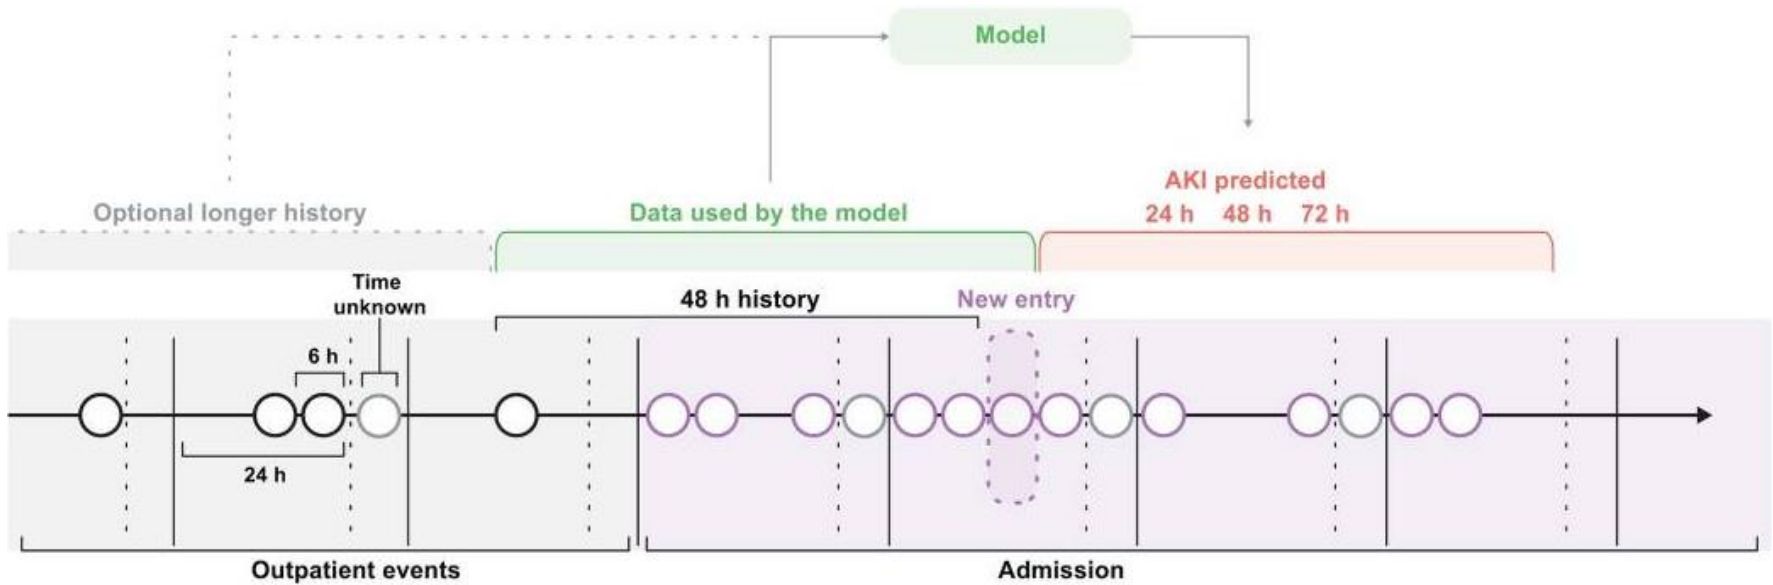

# Auxiliary task



Source: Cortes and Sanchez. IEEE Latin America Transaction (2021)

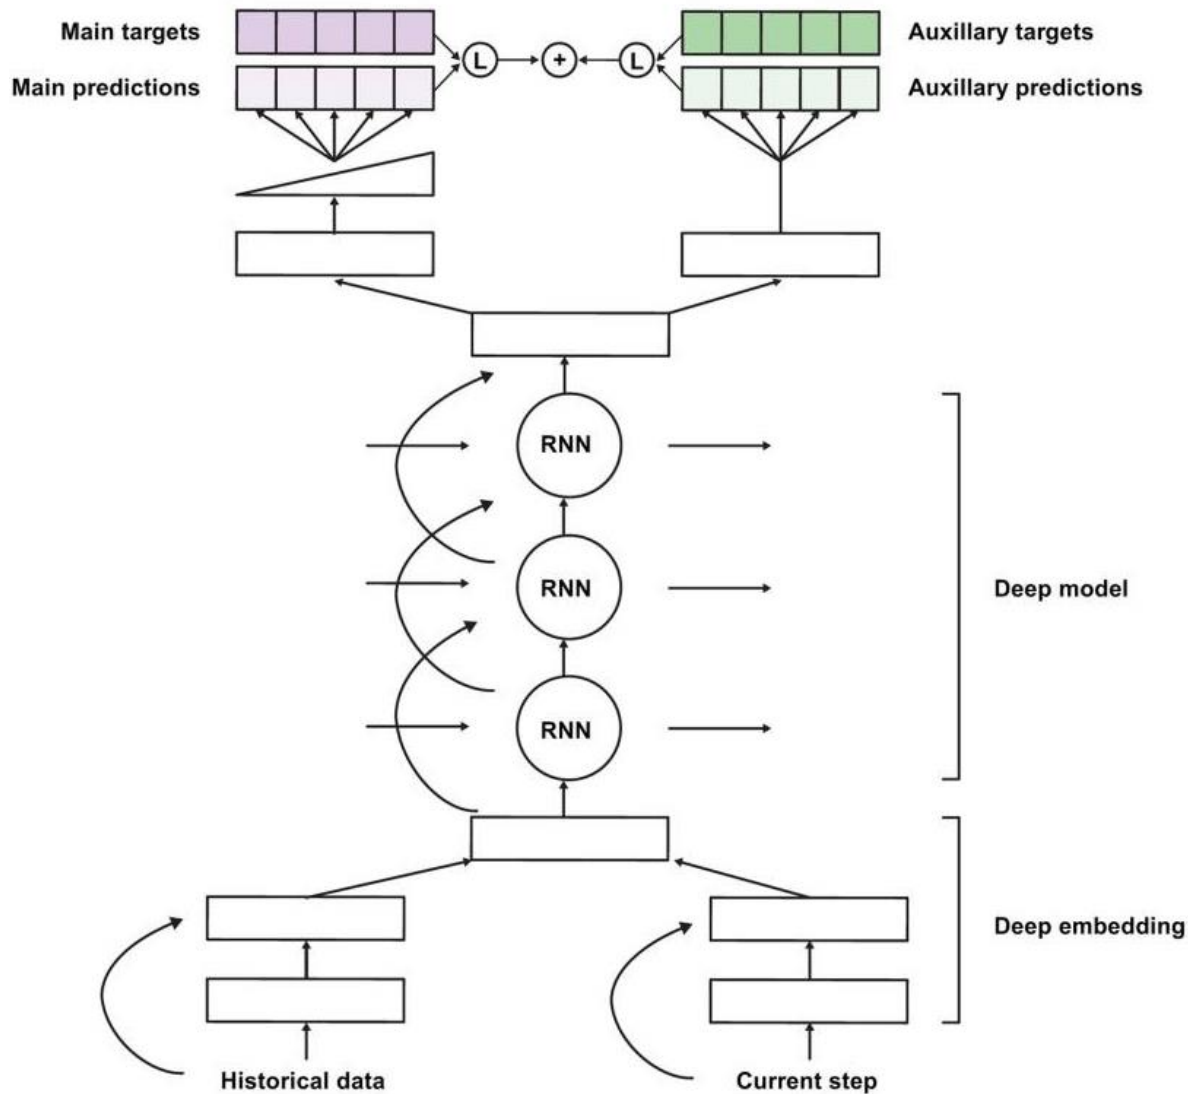
- Good representation should capture multiple characteristics of the data
- Auxiliary task(s) guides the early layers toward good representation

# Acute kidney injury prediction





# Auxiliary tasks





**Main task:** Predict occurrence of AKI within the next 48 hours

**Auxiliary tasks:** Predict maximum values of 7 laboratory features over the next 48 hours

Addition of auxiliary task improve AUC by 3%

# Image segmentation networks

# Segmentation = pixel-level classification




Makropoulos et al. IEEE Trans Med Imaging (2014)

Source: [github.com/kaichoulyc/tgs-salts](https://github.com/kaichoulyc/tgs-salts)


- Output dimension = input dimension!

# Encoder-decoder




# Fully convolutional network

Up-sampling layers



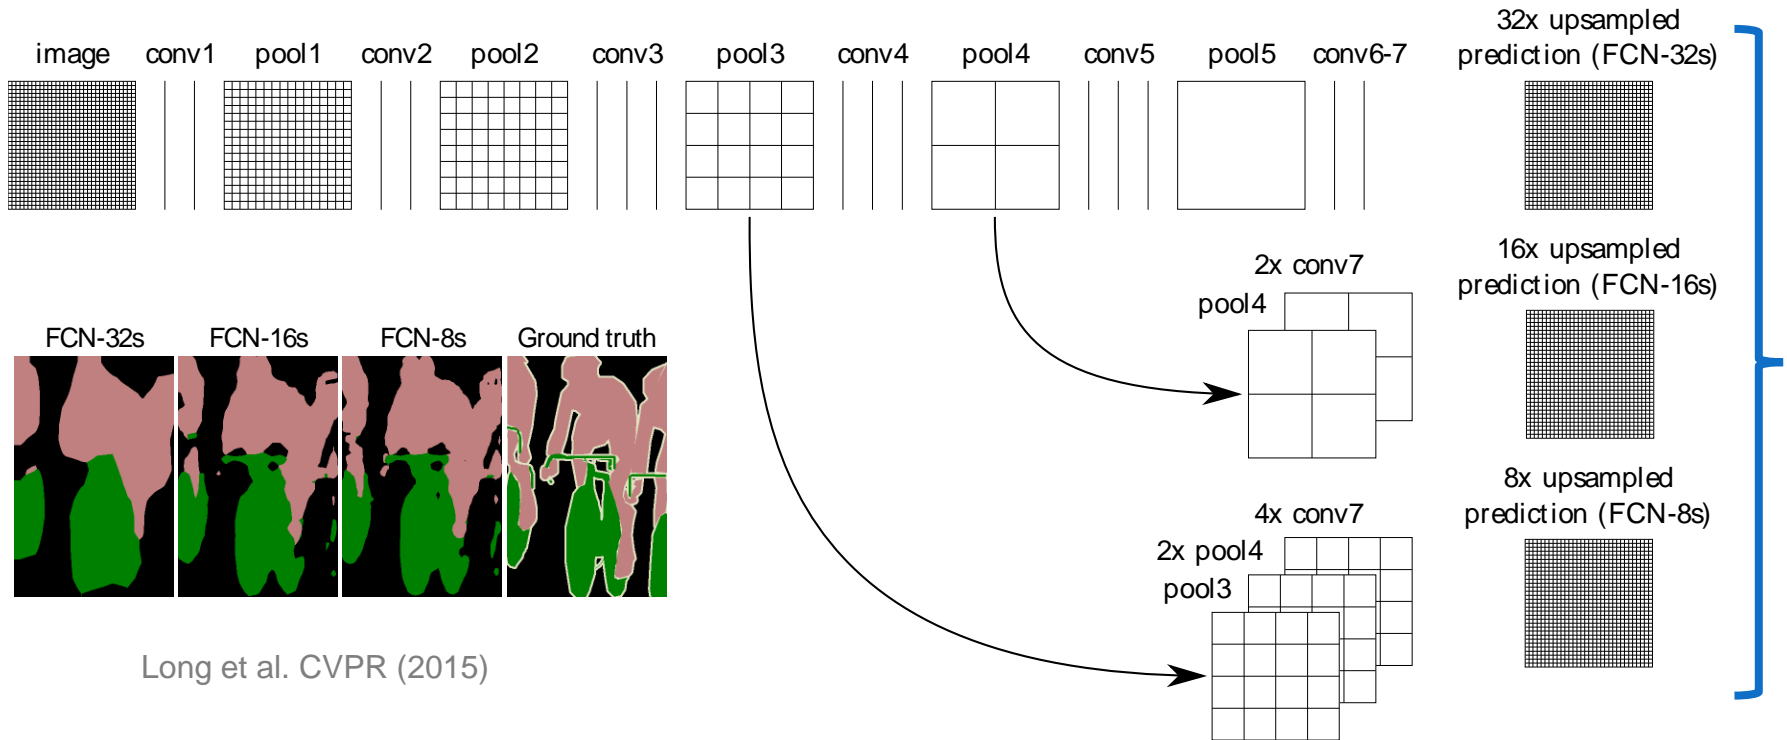
- Using fully connected layers lose spatial and context information

# Up-sampling operation




- Pad input (smaller) image with zeroes
- Apply convolution operation as usual
- Also called “deconvolution” or “transposed convolution”

# Resolution vs feature complexity



- Each successive convolutional layer **increase feature complexity** but **reduce resolution**
- Final outputs = sum of up-sampled outputs from intermediate convolutional layers

# U-Net






# DeepLab

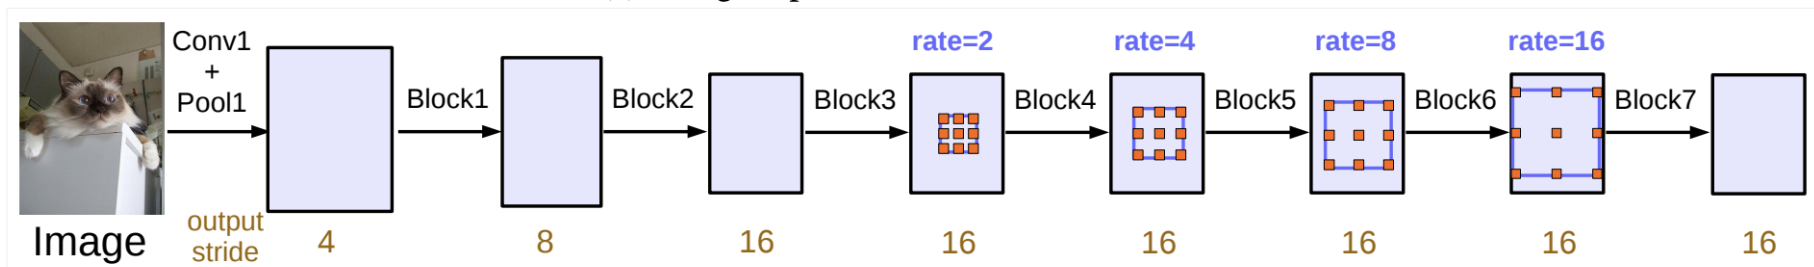


From left: A dog occupying most of the image, a dog occupying a part of it, and a dog occupying very little space (Images obtained from [Unsplash](#)).

Chen et al. CVPR (2017)



(a) Going deeper without atrous convolution.




(b) Going deeper with atrous convolution. Atrous convolution with  $rate > 1$  is applied after block3 when  $output\_stride = 16$ .


Figure 3. Cascaded modules without and with atrous convolution.

- Handle multiple resolutions with **atrous convolution**
  - Same kernel size but cover pixels from broader area

# Limitation of fully convolutional operation



Coarse output from the pixel-wise classifier



MRF/CRF modelling

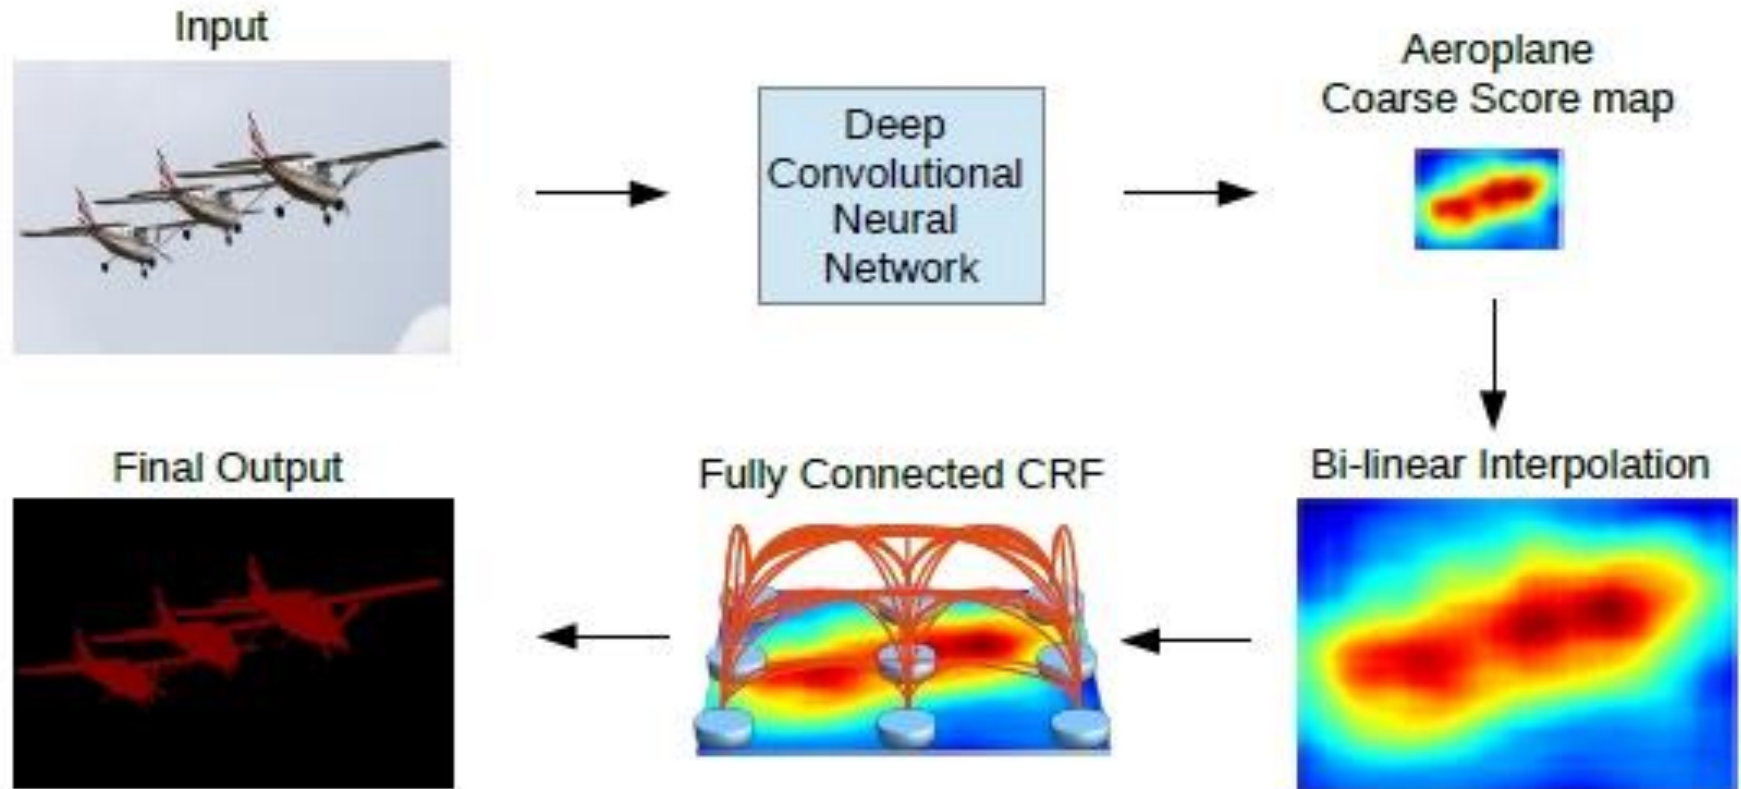


Output after the CRF inference

Source: [towardsdatascience.com](http://towardsdatascience.com)

- Output from fully convolutional network does not take full advantage of information from surrounding pixels
- Additional post-processing with dense layer or conditional/Markov random field

# Conditional random field (CRF)




Chen et al. <http://arxiv.org/pdf/1606.00915> (2016)

- Nearby pixels are conditionally dependent

# Data augmentation


# Small change input → small change in output



Li et al. Neuroimage 183 (2018)

- Slight modifications of input shouldn't change the output by much
  - Same label for classification problem
- ANN should be a smooth function of input

# Real data have variations




Source: [analyticsvidhya.com](https://analyticsvidhya.com)

- ANN should be robust to data noises and variations



# Data augmentation

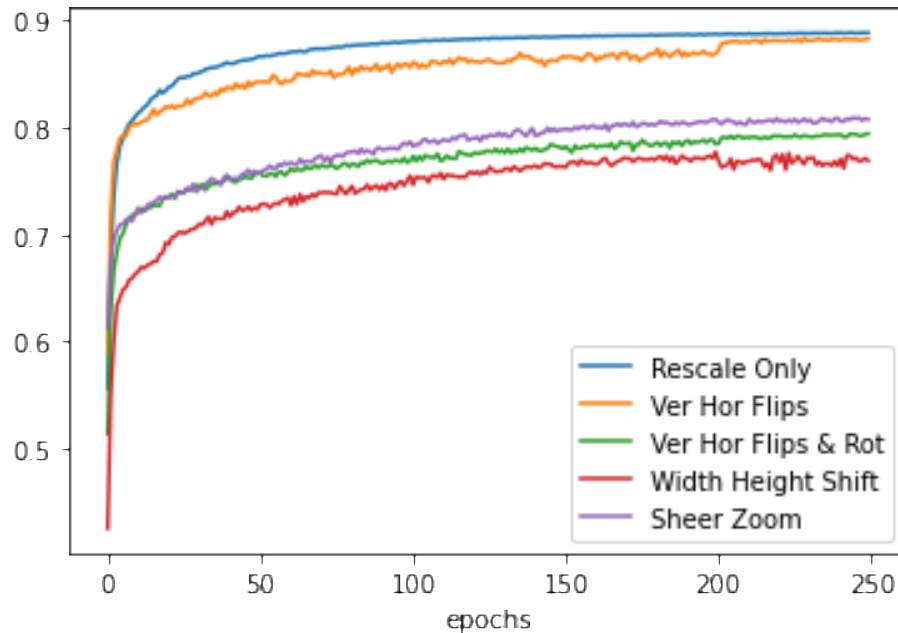


Rama et al. TIPCV (2019)

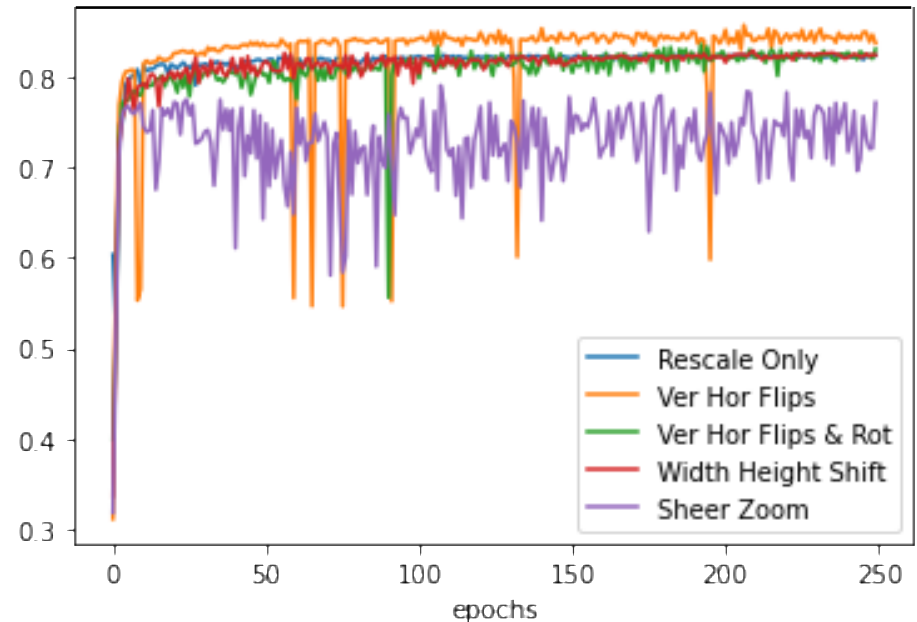
- Randomly perform on each batch of training data
- Help model learn features that are generalized

# Impact of data augmentation

Train accuracy



Validation accuracy



Source: [towardsdatascience.com](https://towardsdatascience.com)

- Too much augmentation → low training accuracy
- Proper augmentation → improved validation
  - Vertical + horizontal flip



# Augmentation in TF

```
tf.keras.preprocessing.image.ImageDataGenerator(  
    featurewise_center=False, samplewise_center=False,  
    featurewise_std_normalization=False, samplewise_std_normalization=False,  
    zca_whitening=False, zca_epsilon=1e-06, rotation_range=0, width_shift_range=0.0,  
    height_shift_range=0.0, brightness_range=None, shear_range=0.0, zoom_range=0.0,  
    channel_shift_range=0.0, fill_mode='nearest', cval=0.0,  
    horizontal_flip=False, vertical_flip=False, rescale=None,  
    preprocessing_function=None, data_format=None, validation_split=0.0, dtype=None  
)
```

```
datagen = ImageDataGenerator(  
    featurewise_center=True,  
    featurewise_std_normalization=True,  
    rotation_range=20,  
    width_shift_range=0.2,  
    height_shift_range=0.2,  
    horizontal_flip=True)  
  
# compute quantities required for featurewise normalization  
# (std, mean, and principal components if ZCA whitening is applied)  
datagen.fit(x_train)  
  
# fits the model on batches with real-time data augmentation:  
model.fit(datagen.flow(x_train, y_train, batch_size=32),  
          steps_per_epoch=len(x_train) / 32, epochs=epochs)
```

Any question?